

IT-641 Deep Learning

#Lab 1

1. Introduction

Machine Learning Pipeline

During this Lab Session we shall revise classification and regression tasks using standard machine learning algorithms. Moreover we shall also try and define a machine learning pipeline that shall help us develop more complex deep learning algorithms in the future:

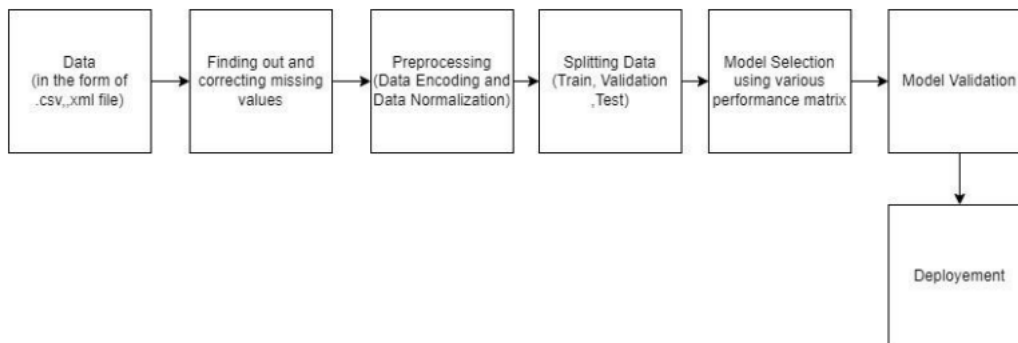


Figure 1(Machine Learning Flowchart)

2. Datasets

1. User dataset This dataset contains information of users from the company's database. It contains information about UserID, Gender, Age, EstimatedSalary, Purchased.
2. Pima Indians Diabetes Database This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. The datasets consist of several medical predictor variables and one target variable, Outcome. Predictor variables include the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.
3. 50_Startups This dataset has data collected from New York, California and Florida about 50 business Startups. The variables used in the dataset are Profit, R&D spending, Administration Spending, and Marketing Spending.

3. Tasks

For each of the above given datasets

1. Load Data and check if the data has missing value
2. Identify which features need to be encoded and encode them
3. Identify which features to normalize and normalize them
4. Identify whether the given task is of classification or regression
5. Split the data into train set (75%) validation set (10%) and test set (15%)
6. Fit the data into 2 models of your choice

REFERENCE CODE - <https://colab.research.google.com/drive/1IEqUTriS66KBbn1V648NooaqlsTB3MvE?usp=sharing>

DATASET 1

User dataset

This dataset contains information of users from the company's database. It contains information about UserID, Gender, Age, EstimatedSalary,Purchased.

1.Loading Required Libraries

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sb
5 sb.set_style("whitegrid")
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import MinMaxScaler
8 from sklearn.preprocessing import StandardScaler,LabelEncoder,LabelBinarizer
9 from sklearn.linear_model import LogisticRegression,SGDClassifier
10 from sklearn.svm import SVC
11 from sklearn.neighbors import KNeighborsClassifier
12 from sklearn.ensemble import RandomForestClassifier
13 from sklearn.metrics import precision_score,recall_score,f1_score,confusion_matrix
14 import warnings
15 warnings.filterwarnings(action = "ignore")
```

2.Loading Data

```
1 data = pd.read_csv("https://raw.githubusercontent.com/Jatansahu/DEEP_LEARNING_ASSIGNMENTS/main/LAB_01/
```

```
1 # Previewing data
2 data.head(8)
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
5	15728773	Male	27	58000	0
6	15598044	Female	27	84000	0
7	15694829	Female	32	150000	1

Looking at the above dataset our target variable is the column "Purchased"

3.Looking for Null values

```
1 print(data.isnull().sum())
```

```
User ID      0
Gender       0
Age          0
EstimatedSalary  0
Purchased    0
dtype: int64
```

There is no null values in our dataset so we will go forward

4.Preprocessing

```
1 data.head()
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

▼ 1.Removing Unnecessary columns

Feature 'User ID' are meaningless when we fit them to our model. Thus we drop these feature.

```
1 data.drop(["User ID"],1,inplace = True)
```

▼ 2.Converting Categorical Variables into their corresponding form

```
1 print(data.dtypes)
```

```
Gender          object
Age             int64
EstimatedSalary int64
Purchased       int64
dtype: object
```

```
1 #encoding the Gender column
2 lb = LabelBinarizer()
3 data['Gender'] = lb.fit_transform(data['Gender'])
```

▼ 3.Scaling Features

In the same way as encoding features we can also scale features manually. Scikit learn as inbuilt scalers that do the same task. Here we shall use standard scaler for our task

```
1 data.describe()
```

	Gender	Age	EstimatedSalary	Purchased
count	400.000000	400.000000	400.000000	400.000000
mean	0.490000	37.655000	69742.500000	0.357500
std	0.500526	10.482877	34096.960282	0.479864
min	0.000000	18.000000	15000.000000	0.000000
25%	0.000000	29.750000	43000.000000	0.000000
50%	0.000000	37.000000	70000.000000	0.000000
75%	1.000000	46.000000	88000.000000	1.000000
max	1.000000	60.000000	150000.000000	1.000000

```
1 # sc = StandardScaler()
2 sc = MinMaxScaler()
3
4 # Fit and transform the data using the scaler
5 # X_scaled = scaler.fit_transform(X)
6 data["EstimatedSalary"] = sc.fit_transform(data["EstimatedSalary"].values.reshape(-1,1))
```

```
1 # sc = StandardScaler()
2 sc = MinMaxScaler()
3 data["Age"] = sc.fit_transform(data["Age"].values.reshape(-1,1))
```

5.Basic EDA

1. Gathering some info about data

```
1 data.describe().T
```


	count	mean	std	min	25%	50%	75%	max
Gender	400.0	0.490000	0.500526	0.0	0.000000	0.000000	1.000000	1.0
Age	400.0	0.467976	0.249592	0.0	0.279762	0.452381	0.666667	1.0
EstimatedSalary	400.0	0.405500	0.252570	0.0	0.207407	0.407407	0.540741	1.0
Purchased	400.0	0.357500	0.479864	0.0	0.000000	0.000000	1.000000	1.0

```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Gender                400 non-null   int64   
1   Age                   400 non-null   float64  
2   EstimatedSalary       400 non-null   float64  
3   Purchased             400 non-null   int64   
dtypes: float64(2), int64(2)
memory usage: 12.6 KB
```

2. Correlation plot

```
1 plt.figure(figsize = (25,10))
2 sb.heatmap(data.corr(),annot = True);
```



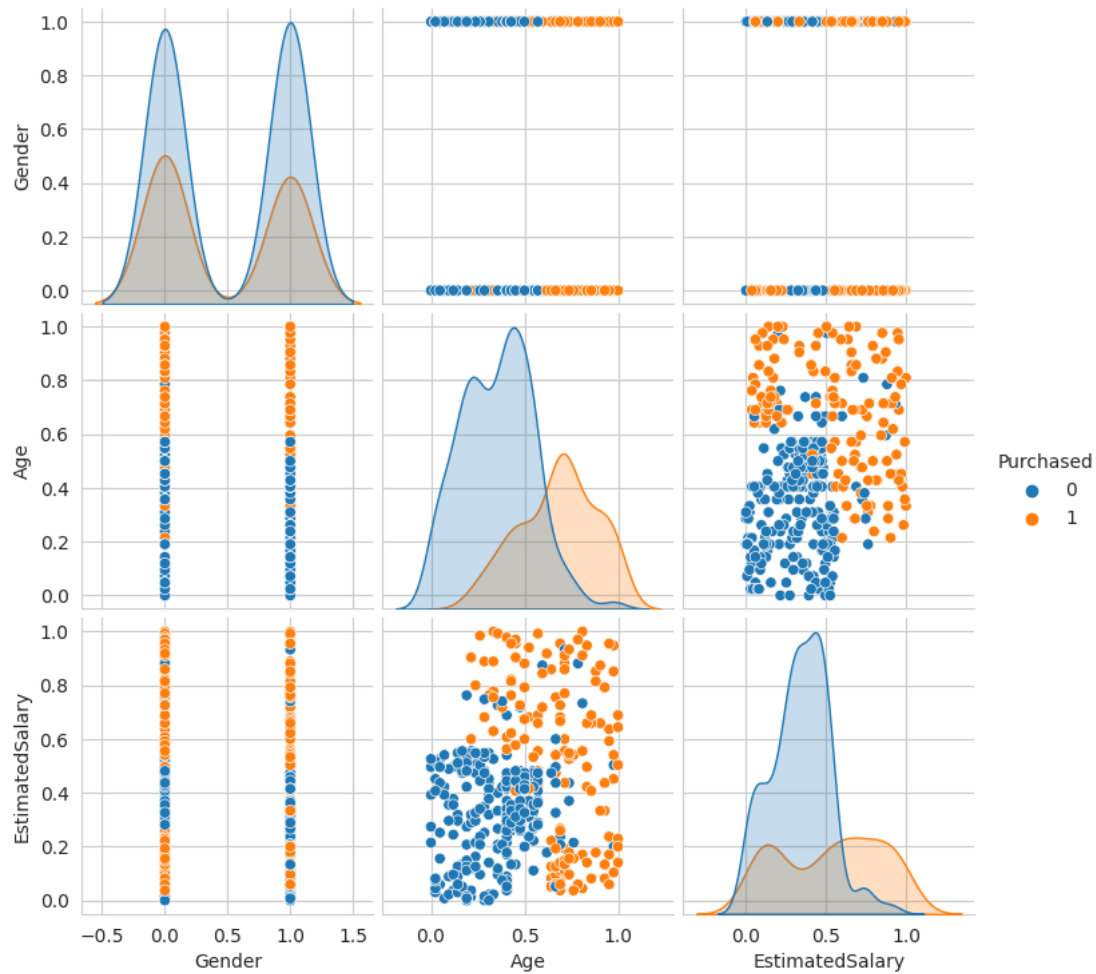
OBSERVATION:

We clearly see that attribute **Age** and **Purchased** attribute have a correlation of 0.62 suggests a moderately strong positive correlation. It implies that as age increases, the purchases tend to increase as well, but not necessarily in a perfectly linear fashion

3. Pairplot



```
1 plt.figure(figsize = (25,25))
2 sb.pairplot(data,hue = "Purchased");
```

<Figure size 2500x2500 with 0 Axes>



▼ 4.You can plot using pandas too.

```
1 data.hist(figsize = (15,10),color = 'blue');
```



6. Splitting the dataset

Split the data into train set (75%) validation set (10%) and test set (15%)



1 data

	Gender	Age	EstimatedSalary	Purchased
0	1	0.023810	0.029630	0
1	1	0.404762	0.037037	0
2	0	0.190476	0.207407	0
3	0	0.214286	0.311111	0
4	1	0.023810	0.451852	0
...
395	0	0.666667	0.192593	1
396	1	0.785714	0.059259	1
397	0	0.761905	0.037037	1
398	1	0.428571	0.133333	0
399	0	0.738095	0.155556	1

400 rows × 4 columns

```
1 x = data.iloc[:, :3]
2 y = data['Purchased']
```

5. Split the data into train set (75%) validation set (10%) and test set (15%)

```
1 x_train, x_part, y_train, y_part = train_test_split(x, y, test_size = 0.25, random_state = 42)
2 x_test, x_valid, y_test, y_valid = train_test_split(x_part, y_part, test_size = 0.4, random_state = 42)

1 print(x_train.shape, x_test.shape, x_valid.shape)
2 print(y_train.shape, y_test.shape, y_valid.shape)

(300, 3) (60, 3) (40, 3)
(300,) (60,) (40,)
```

7. Model Selection

Before we fit our data into our model we need to define some metrics with the help of which we can select the best fitting model

As our current task is classification we shall create a function that evaluates our model based on precision score, recall score and F1-score

```
1 def evaluate(model, model_name, x_train = x_train, y_train = y_train, x_test = x_test, y_test = y_test, x_val
2   print(f"Model performance for {model_name}")
3   y_train_pred = model.predict(x_train)
4   y_test_pred = model.predict(x_test)
5   y_valid_pred = model.predict(x_valid)
```

```

6
7 #confusion matrix
8 plt.figure(figsize = (10,10))
9 sb.heatmap(confusion_matrix(y_train,y_train_pred),annot = True)
10 plt.title('Confusion Matrix')
11 plt.show()
12
13 #precision score
14 precision_score_train = precision_score(y_train,y_train_pred)
15 precision_score_test = precision_score(y_test,y_test_pred)
16 precision_score_valid = precision_score(y_valid,y_valid_pred)
17
18 #recallscore
19 recall_score_train = recall_score(y_train,y_train_pred)
20 recall_score_test = recall_score(y_test,y_test_pred)
21 recall_score_valid = recall_score(y_valid,y_valid_pred)
22
23 #f1 score
24 f1_score_train = f1_score(y_train,y_train_pred)
25 f1_score_test = f1_score(y_test,y_test_pred)
26 f1_score_valid = f1_score(y_valid,y_valid_pred)
27
28 print("Precision Score Train:",precision_score_train)
29 print("Precision Score Test:",precision_score_test)
30 print("Precision Score Validation",precision_score_valid)
31
32 print("recall Score Train:",recall_score_train)
33 print("recal Score Test:",recall_score_test)
34 print("recall Score Validation",recall_score_valid)
35
36 print("f1 Score Train:",f1_score_train)
37 print("f1 Score Test:",f1_score_test)
38 print("f1 Score Validation",f1_score_valid)
39
40
41
42 return precision_score_train,precision_score_test,precision_score_valid,recall_score_train,recall_sc
43

```

```

1 clf1 = LogisticRegression()
2 clf1.fit(x_train,y_train)

```

```

▼ LogisticRegression
LogisticRegression()


```

```

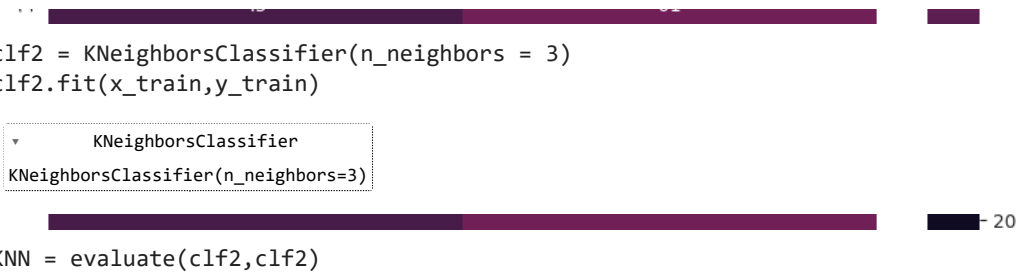
1 LR = evaluate(clf1,clf1)

```


Model performance forLogisticRegression()



2.KNN



Model performance for KNeighborsClassifier(n_neighbors=3)



3.SVM

```
1 clf3 = SVC(kernel = "linear")  
2 clf3.fit(x_train,y_train)
```

SVC
SVC(kernel='linear')

```
1 svc = evaluate(clf3,clf3)
```

Model performance for SVC(kernel='linear')

Confusion Matrix



4 Random Forest



```
1 clf4 = RandomForestClassifier(n_estimators=100, random_state=42)
2 clf4.fit(x_train,y_train)
```


```
RandomForestClassifier
RandomForestClassifier(random_state=42)
```



```
1 rf = evaluate(clf4,clf4)
```

Model performance for RandomForestClassifier(random_state=42)

Confusion Matrix



```
Precision Score Train: 0.9906542056074766
Precision Score Test: 0.9444444444444444
Precision Score Validation 0.75
recall Score Train: 1.0
recall Score Test: 0.7727272727272727
recall Score Validation 1.0
f1 Score Train: 0.9953051643192489
f1 Score Test: 0.85
f1 Score Validation 0.8571428571428571
```

```
1 model_performance = [{"Logistic Regression", LR[0], LR[1], LR[3], LR[4], LR[5], LR[6], LR[7], LR[8]],
2                       ["Knearest Neighbors", KNN[0], KNN[1], KNN[3], KNN[4], KNN[5], KNN[6], KNN[7], KNN[8]],
3                       ["Support Vector Machine", svc[0], svc[1], svc[3], svc[4], svc[5], svc[6], svc[7], svc[8]],
4                       ["Random Forest", rf[0], rf[1], rf[3], rf[4], rf[5], rf[6], rf[7], rf[8]]
5                       ]
```

```
1 model_performance = pd.DataFrame(model_performance,columns = ["Precision Score Train","Precision Score
```

```
1 model_performance
```

		Precision Score Train	Precision Score Test	Precision Score Validation	Recall Score Train	Recall Score Test	Recall Score Validation	F1 Score Train	F1 Score Test	F1 Score Validation
0	Logistic Regression		0.824324	0.937500	0.575472	0.681818	0.666667	0.677778	0.789474	0.800000
1	Knearest Neighbors		0.867257	0.950000	0.924528	0.863636	0.933333	0.894977	0.904762	0.848485
2	Support Vector Machine		0.828947	0.937500	0.594340	0.681818	0.666667	0.692308	0.789474	0.800000
3	Random Forest		0.990654	0.944444	1.000000	0.772727	1.000000	0.995305	0.850000	0.857143

⬆ ⬇ ⬆

DATASET - 2

Pima Indians Diabetes Database

This dataset is originally from the National Institute of Diabetes and Digestive and Kidney Diseases. The objective of the dataset is to diagnostically predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. The datasets consist of several medical predictor variables and one target variable, Outcome. Predictor variables include the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

1.Loading Required Libraries

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sb
5 sb.set_style("whitegrid")
6 from sklearn.model_selection import train_test_split
7 from sklearn.preprocessing import StandardScaler,LabelEncoder,LabelBinarizer
8 from sklearn.linear_model import LogisticRegression,SGDClassifier
9 from sklearn.linear_model import LogisticRegression,SGDClassifier
10 from sklearn.ensemble import RandomForestClassifier
11 from sklearn.svm import SVC
12 from sklearn.neighbors import KNeighborsClassifier
13 from sklearn.metrics import precision_score,recall_score,f1_score,confusion_matrix
14 import warnings
15 warnings.filterwarnings(action = "ignore")
```

2.Loading Data

```
1 data = pd.read_csv("https://raw.githubusercontent.com/Jatansahu/DEEP_LEARNING_ASSIGNMENTS/main/LAB_01/
```

```
1 # Previewing data
2 data.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Looking at the above dataset our target variable is the column "Outcome"

3.Looking for Null values

```
1 print(data.isnull().sum())
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness    0
Insulin           0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

There is no missing or null values in the dataset

4.Preprocessing

```
1 data.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

1.Removing Unnecessary columns

```
1 data.columns
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

Based on general domain knowledge, some features may be considered more directly related to diabetes risk than others. In many cases, "Pregnancies" might not be directly related to diabetes risk but could have an indirect impact through other factors. It's important to conduct a thorough analysis, such as feature importance from a machine learning model, to determine the relative importance of each feature in predicting diabetes for a given dataset.

2.Converting Categorical Variables into their corresponding form

```
1 print(data.dtypes)
```

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness    int64
Insulin           int64
BMI              float64
DiabetesPedigreeFunction float64
Age              int64
Outcome          int64
dtype: object
```

There is no categorical variables in the dataset.


3.Scaling Features

StandardScaler: This scaler assumes that the data follows a Gaussian (normal) distribution

MinMaxScaler: This scaler is more appropriate when your data doesn't follow a normal distribution or when you have features with significantly different scales.

```
1 sc = StandardScaler()
2 data["Pregnancies"] = sc.fit_transform(data["Pregnancies"].values.reshape(-1,1))
```

```
1 sb.distplot(data["Glucose"]);
2 print("Mean", np.round(np.mean(data["Glucose"]),0))
```



```
1 sc = StandardScaler()
2 data["Glucose"] = sc.fit_transform(data["Glucose"].values.reshape(-1,1))
```

```
1 sc = StandardScaler()
2 data["BloodPressure"] = sc.fit_transform(data["BloodPressure"].values.reshape(-1,1))
```

```
1 sc = StandardScaler()
2 data["Insulin"] = sc.fit_transform(data["Insulin"].values.reshape(-1,1))
```

```
1 sc = StandardScaler()
2 data["BMI"] = sc.fit_transform(data["BMI"].values.reshape(-1,1))
```

```
1 sc = StandardScaler()
2 data["DiabetesPedigreeFunction"] = sc.fit_transform(data["DiabetesPedigreeFunction"].values.reshape(-1,1))
```

```
1 sc = StandardScaler()
2 data["Age"] = sc.fit_transform(data["Age"].values.reshape(-1,1))
```

```
1 sc = StandardScaler()
2 data["SkinThickness"] = sc.fit_transform(data["SkinThickness"].values.reshape(-1,1))
```

5.Basic EDA

▼ 1.Gathering some info about data

```
1 data.describe().T
```

	count	mean	std	min	25%	50%	75%	max
Pregnancies	768.0	-6.476301e-17	1.000652	-1.141852	-0.844885	-0.250952	0.639947	3.906578
Glucose	768.0	-9.251859e-18	1.000652	-3.783654	-0.685236	-0.121888	0.605771	2.444478
BloodPressure	768.0	1.503427e-17	1.000652	-3.572597	-0.367337	0.149641	0.563223	2.734528
SkinThickness	768.0	1.006140e-16	1.000652	-1.288212	-1.288212	0.154533	0.719086	4.921866
Insulin	768.0	-3.006854e-17	1.000652	-0.692891	-0.692891	-0.428062	0.412008	6.652839
BMI	768.0	2.590520e-16	1.000652	-4.060474	-0.595578	0.000942	0.584771	4.455807
DiabetesPedigreeFunction	768.0	2.451743e-16	1.000652	-1.189553	-0.688969	-0.300128	0.466227	5.883565
Age	768.0	1.021225e-16	1.000652	-1.041540	-0.786286	-0.360847	0.660206	1.062716




```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   float64
1   Glucose                768 non-null   float64
2   BloodPressure          768 non-null   float64
3   SkinThickness          768 non-null   float64
4   Insulin                768 non-null   float64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   float64
8   Outcome                768 non-null   int64
dtypes: float64(8), int64(1)
memory usage: 54.1 KB
```

2. Correlation plot

```
1 plt.figure(figsize = (25,10))
2 sb.heatmap(data.corr(),annot = True);
```



In general, a common approach is to set a correlation threshold (often a positive value) and keep features with correlations above that threshold. Common threshold values can range from 0.1 to 0.3

In this case I have decided the threshold value of 0.1

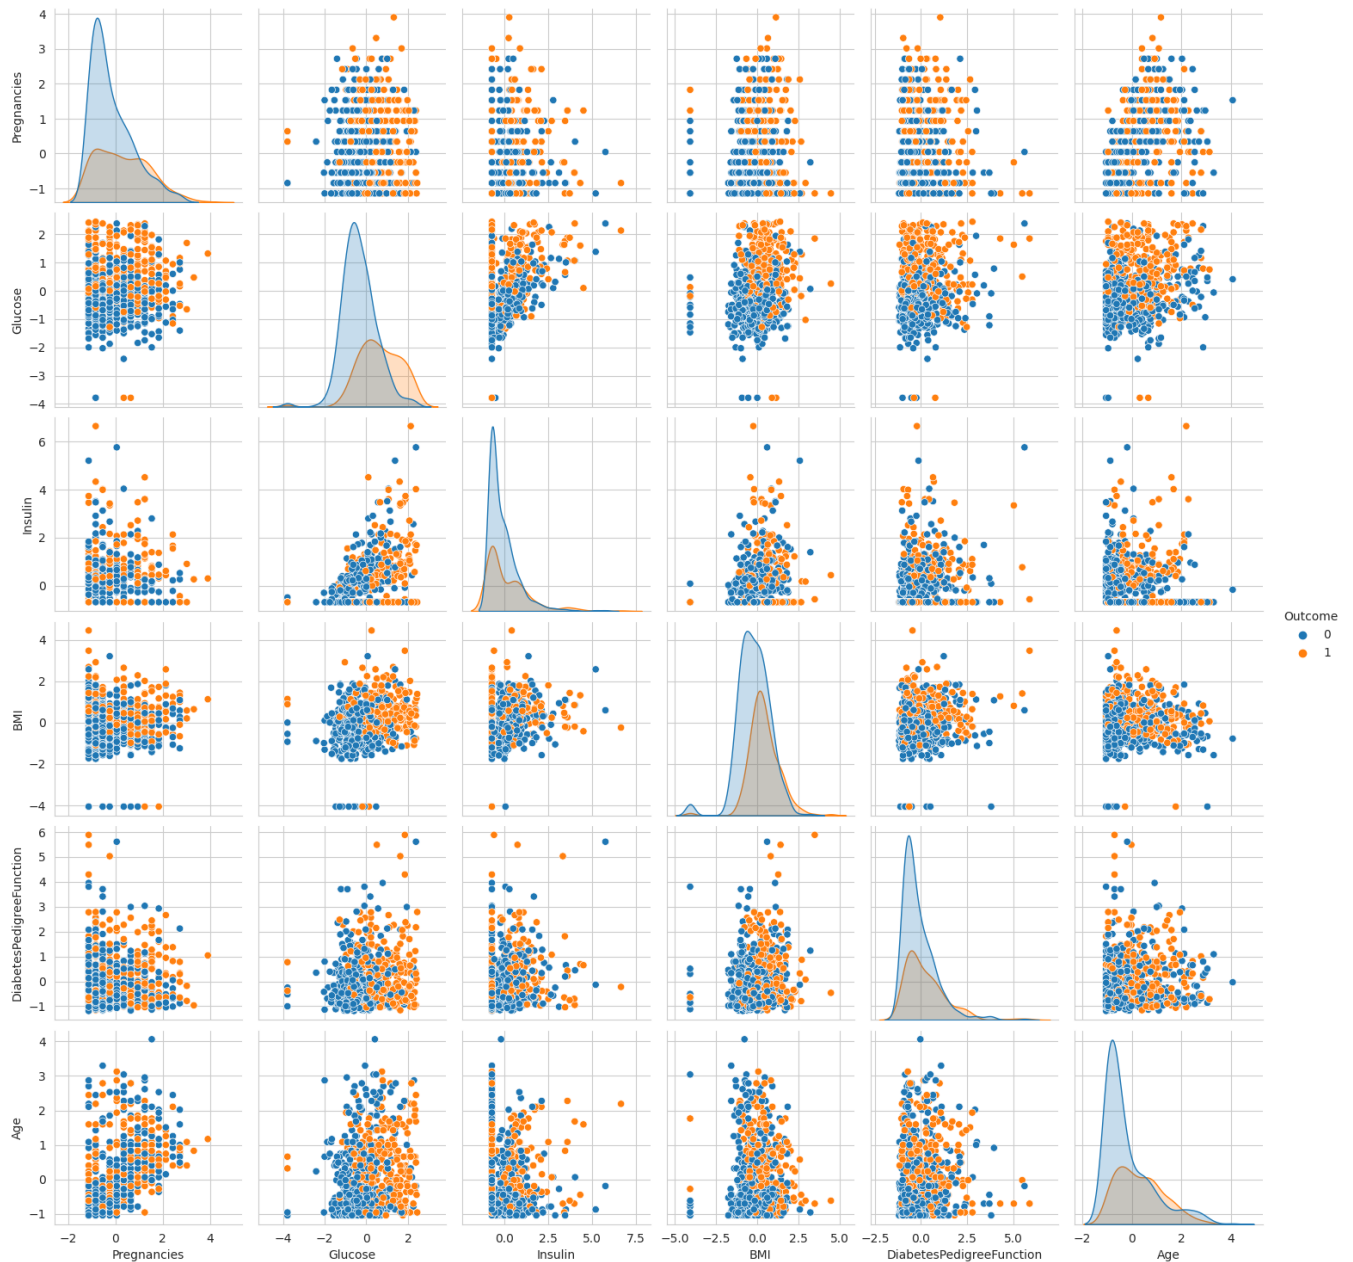
From above correlation chart we are rejecting BloodPressure and Skinthickness feature.

```
1 data.drop(["BloodPressure","SkinThickness"],1,inplace = True)
```

▼ 3. Pairplot

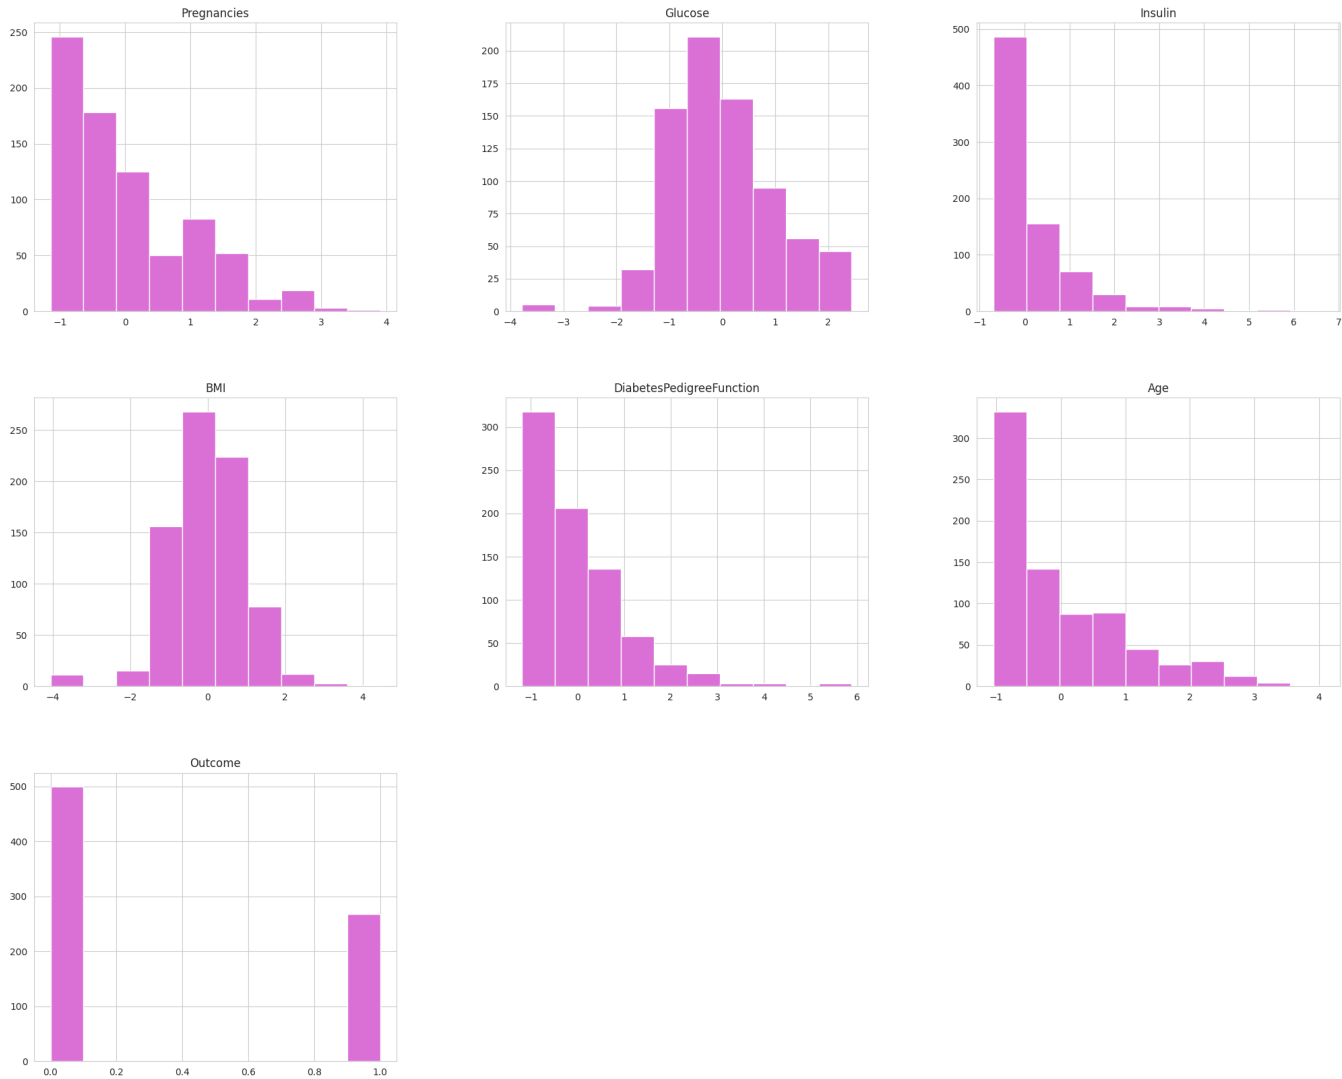
```
1 plt.figure(figsize = (25,25))
2 sb.pairplot(data,hue = "Outcome");
```

<Figure size 2500x2500 with 0 Axes>



▼ 4. You can plot using pandas too..

```
1 data.hist(figsize = (25,20),color = 'orchid');
2
```



6.Splitting the dataset

1 data

	Pregnancies	Glucose	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	0.639947	0.848324	-0.692891	0.204013	0.468492	1.425995	1
1	-0.844885	-1.123396	-0.692891	-0.684422	-0.365061	-0.190672	0
2	1.233880	1.943724	-0.692891	-1.103255	0.604397	-0.105584	1
3	-0.844885	-0.998208	0.123302	-0.494043	-0.920763	-1.041549	0
4	-1.141852	0.504055	0.765836	1.409746	5.484909	-0.020496	1
...
763	1.827813	-0.622642	0.870031	0.115169	-0.908682	2.532136	0
764	-0.547919	0.034598	-0.692891	0.610154	-0.398282	-0.531023	0
765	0.342981	0.003301	0.279594	-0.735190	-0.685193	-0.275760	0
766	-0.844885	0.159787	-0.692891	-0.240205	-0.371101	1.170732	1
767	-0.844885	-0.873019	-0.692891	-0.202129	-0.473785	-0.871374	0

768 rows × 7 columns


```
1 x = data.iloc[:, :6]
2 y = data['Outcome']
```

Split the data into train set (75%) validation set (10%) and test set (15%)

```
1 x_train,x_part,y_train,y_part = train_test_split(x,y,test_size = 0.25,random_state = 42)
2 x_test,x_valid,y_test,y_valid = train_test_split(x_part,y_part,test_size = 0.4,random_state = 42)

1 print(x_train.shape,x_test.shape,x_valid.shape)
2 print(y_train.shape,y_test.shape,y_valid.shape)

(576, 6) (115, 6) (77, 6)
(576,) (115,) (77,)
```

▼ 7.Model Selection

Before we fit our data into our model we need to define some metrics with the help of which we can select the best fitting model

As our current task is classification we shall create a function that evaluates our model based on precision score,recall score and F1-score

```
1 def evaluate(model,model_name,x_train = x_train,y_train = y_train,x_test = x_test,y_test = y_test,x_va
2   print(f"Model performance for{model_name}")
3   y_train_pred = model.predict(x_train)
4   y_test_pred = model.predict(x_test)
5   y_valid_pred = model.predict(x_valid)
6
7   #confusion matrix
8   plt.figure(figsize = (10,10))
9   sb.heatmap(confusion_matrix(y_train,y_train_pred),annot = True)
10  plt.title('Confusion Matrix')
11  plt.show()
12
13  #precision score
14  precision_score_train = precision_score(y_train,y_train_pred)
15  precision_score_test = precision_score(y_test,y_test_pred)
16  precision_score_valid = precision_score(y_valid,y_valid_pred)
17
18  #recallscore
19  recall_score_train = recall_score(y_train,y_train_pred)
20  recall_score_test = recall_score(y_test,y_test_pred)
21  recall_score_valid = recall_score(y_valid,y_valid_pred)
22
23  #f1 score
24  f1_score_train = f1_score(y_train,y_train_pred)
25  f1_score_test = f1_score(y_test,y_test_pred)
26  f1_score_valid = f1_score(y_valid,y_valid_pred)
27
28  print("Precision Score Train:",precision_score_train)
29  print("Precision Score Test:",precision_score_test)
30  print("Precision Score Validation",precision_score_valid)
31
32  print("recall Score Train:",recall_score_train)
33  print("recal Score Test:",recall_score_test)
34  print("recall Score Validation",recall_score_valid)
35
36  print("f1 Score Train:",f1_score_train)
37  print("f1 Score Test:",f1_score_test)
38  print("f1 Score Validation",f1_score_valid)
39
40
41
42  return precision_score_train,precision_score_test,precision_score_valid,recall_score_train,recall_sc
43
```

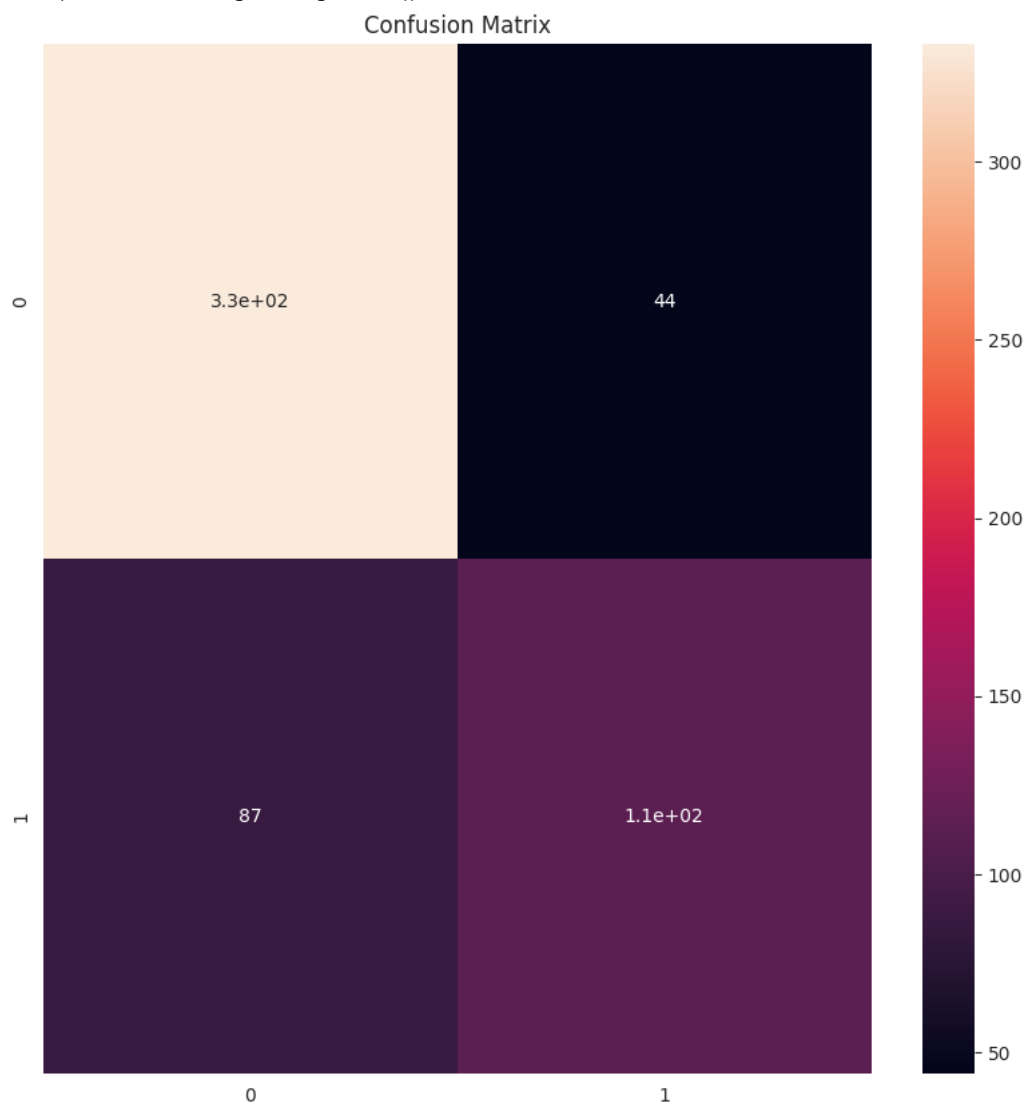
▼ 1. Logistic Regression

```
1 clf1 = LogisticRegression()
2 clf1.fit(x_train,y_train)
```

```
▼ LogisticRegression
LogisticRegression()
```

```
1 LR = evaluate(clf1,clf1)
```

Model performance for LogisticRegression()



```
Precision Score Train: 0.717948717948718
Precision Score Test: 0.6285714285714286
Precision Score Validation 0.6363636363636364
recall Score Train: 0.5628140703517588
recal Score Test: 0.55
recall Score Validation 0.7241379310344828
f1 Score Train: 0.6309859154929578
f1 Score Test: 0.5866666666666667
f1 Score Validation 0.6774193548387097
```


▼ 2.KNN

```
1 clf2 = KNeighborsClassifier(n_neighbors = 3)
2 clf2.fit(x_train,y_train)
```

```
▼ KNeighborsClassifier
KNeighborsClassifier(n_neighbors=3)
```

```
1 KNN = evaluate(clf2,clf2)
```

Model performance forKNeighborsClassifier(n_neighbors=3)



```
Precision Score Train: 0.798941798941799
Precision Score Test: 0.5
Precision Score Validation 0.6551724137931034
recall Score Train: 0.7587939698492462
recall Score Test: 0.475
recall Score Validation 0.6551724137931034
f1 Score Train: 0.7783505154639175
f1 Score Test: 0.48717948717948717
f1 Score Validation 0.6551724137931034
```


▼ 3.SVM

```
1 clf3 = SVC(kernel = "linear")
2 clf3.fit(x_train,y_train)
```

```
SVC
SVC(kernel='linear')
```

```
1 svc = evaluate(clf3,clf3)
```


Model performance for SVC(kernel='linear')



1 svc = evaluate(clf3,clf3)

Model performance for SVC(kernel='linear')

Confusion Matrix



```
1 model_performance = ["Logistic Regression", LR[0], LR[1], LR[3], LR[4], LR[5], LR[6], LR[7], LR[8]],
2                       ["Knearest Neighbors", KNN[0], KNN[1], KNN[3], KNN[4], KNN[5], KNN[6], KNN[7], KNN[8]],
3                       ["Support Vector Machine", svc[0], svc[1], svc[3], svc[4], svc[5], svc[6], svc[7], svc[8]]
```

```
1 model_performance = pd.DataFrame(model_performance, columns = ["Precision Score Train", "Precision Score
```

```
1 model_performance
```

	Precision Score Train	Precision Score Test	Precision Score Validation	Recall Score Train	Recall Score Test	Recall Score Validation	F1 Score Train	F1 Score Test	F1 Score Validation
0	Logistic Regression	0.717949	0.628571	0.562814	0.550	0.724138	0.630986	0.586667	0.677419
1	Knearest Neighbors	0.798942	0.500000	0.758794	0.475	0.655172	0.778351	0.487179	0.655172
2	Support Vector Machine	0.717949	0.628571	0.562814	0.550	0.724138	0.630986	0.586667	0.677419

DATASET03 - 50_Startups

50_Startups This dataset has data collected from New York, California and Florida about 50 business Startups. The variables used in the dataset are Profit, R&D spending, Administration Spending, and Marketing Spending.

```
Precision Score Validation: 0.6285714285714286
```

1.Loading Required Libraries

```
f1 Score Test: 0.5866666666666667
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sb
5 from sklearn.preprocessing import LabelEncoder, StandardScaler
6 from sklearn.model_selection import train_test_split, GridSearchCV
7 import warnings
8 warnings.filterwarnings(action = 'ignore')
9 from sklearn.preprocessing import LabelEncoder, StandardScaler
10 from sklearn.linear_model import Lasso, LinearRegression, ElasticNet, Ridge
11 from sklearn.neighbors import KNeighborsRegressor
12 from sklearn.tree import DecisionTreeRegressor
13 from sklearn.ensemble import RandomForestRegressor, AdaBoostRegressor
14 from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, cross_val_predict
15 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
16 import xgboost
17 # import optuna
```

2. Load Dataset

```
1 data = pd.read_csv("https://raw.githubusercontent.com/Jatansahu/DEEP_LEARNING_ASSIGNMENTS/main/LAB_01/
```

```
1 # Previewing data
2 data.head()
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

Looking at the above dataset our target variable is the column "Profit"

3.Looking for Null values

```
1 print(data.isnull().sum())

R&D Spend      0
Administration 0
Marketing Spend 0
State           0
Profit          0
dtype: int64
```

Double-click (or enter) to edit

4.Preprocessing

1. Removing Unnecessary columns

```
1 data['State'].unique()

array(['New York', 'California', 'Florida'], dtype=object)
```

As of now, we don't know which column is not much related

2.Converting Categorical Variables into their corresponding form

```
1 print(data.dtypes)

R&D Spend      float64
Administration float64
Marketing Spend float64
State           object
Profit          float64
dtype: object

1 #encoding Embarked column
2 le = LabelEncoder()
3 data['State'] = le.fit_transform(data["State"])

1 data.head()
```

3. Scaling Features

In the same way as encoding features we can also scale features manually. Scikit learn has inbuilt scalers that do the same task. Here we shall use standard scaler for our task

```
4  142107 34      91391 77      366168 42      1  166187 94
1  sc = StandardScaler()
2  data["R&D Spend"] = sc.fit_transform(data["R&D Spend"].values.reshape(-1,1))
3  data["Administration"] = sc.fit_transform(data["Administration"].values.reshape(-1,1))
4  data["Marketing Spend"] = sc.fit_transform(data["Marketing Spend"].values.reshape(-1,1))
5  # data["Profit"] = sc.fit_transform(data["Profit"].values.reshape(-1,1))
```

There's an important consideration when it comes to interpretation. If we scale the target variable during preprocessing (for example, using MinMaxScaler to scale it to a specific range), we'll need to remember that any predictions made by the model will be in the scaled range. If we need to interpret the predictions in the original units (e.g., dollars for profit), we'll have to reverse the scaling transformation to get the predictions in the original scale."

```
1  profit_data = data[["Profit"]] # Extracting the "Profit" column as a separate DataFrame
2  scaler = StandardScaler()
3  scaled_profit = scaler.fit_transform(profit_data.values.reshape(-1,1))
4  # Converting the scaled profit back to a pandas Series (if needed)
5  # scaled_profit_series = pd.Series(scaled_profit[:, 0], name="Scaled_Profit")

1  # # Get the mean and standard deviation from the scaler
2  # mean_profit = scaler.mean_[0]
3  # std_dev_profit = scaler.scale_[0]
4
5  # scaled_prediction = 2.01120333
6
7  # # Reverse the scaling to get the prediction in the original units
8  # original_prediction = (scaled_prediction * std_dev_profit) + mean_profit
9
10 # print("Original prediction in dollars:", original_prediction)
11

1  # Dropping profit column from dataset
2  data.drop(["Profit"],1,inplace = True)

1  data['scaled_profit'] = scaled_profit

1  data.head()
```

	R&D Spend	Administration	Marketing Spend	State	scaled_profit
0	2.016411	0.560753	2.153943	2	2.011203
1	1.955860	1.082807	1.923600	0	1.999430
2	1.754364	-0.728257	1.626528	1	1.980842
3	1.554784	-0.096365	1.422210	2	1.776627
4	1.504937	-1.079919	1.281528	1	1.357740

5. Basic EDA

1. Gathering some info about data

```
1  data.describe().T
```





```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   R&D Spend              50 non-null     float64
1   Administration         50 non-null     float64
2   Marketing Spend        50 non-null     float64
3   State                  50 non-null     int64
4   scaled_profit          50 non-null     float64
dtypes: float64(4), int64(1)
memory usage: 2.1 KB
```

- 2. Correlation plot

```
1 plt.figure(figsize = (25,10))
2 sb.heatmap(data.corr(),annot = True);
```




We will not take state as a feature in our data preprocessing part

▼ 3.Pairplot


```
1 plt.figure(figsize = (25,25))
2 sb.pairplot(data,hue = "scaled_profit");
```


<Figure size 2500x2500 with 0 Axes>



▼ 4.You can plot using pandas too..

```
1 data.hist(figsize = (25,20),color = 'orchid');
```



6.Splitting the dataset

Split the data into train set (75%) validation set (10%) and test set (15%)



1 data.head()

	R&D Spend	Administration	Marketing Spend	State	scaled_profit
0	2.016411	0.560753	2.153943	2	2.011203
1	1.955860	1.082807	1.923600	0	1.999430
2	1.754364	-0.728257	1.626528	1	1.980842
3	1.554784	-0.096365	1.422210	2	1.776627
4	1.504937	-1.079919	1.281528	1	1.357740

```
1 x = data.iloc[:, :3]
2 y = data['scaled_profit']
```

1 x.head()

	R&D Spend	Administration	Marketing Spend
0	2.016411	0.560753	2.153943
1	1.955860	1.082807	1.923600
2	1.754364	-0.728257	1.626528
3	1.554784	-0.096365	1.422210
4	1.504937	-1.079919	1.281528

```
1 x_train,x_part,y_train,y_part = train_test_split(x,y,test_size = 0.25,random_state = 42)
2 x_test,x_valid,y_test,y_valid = train_test_split(x_part,y_part,test_size = 0.4,random_state = 42)
```

```
1 print(x_train.shape,x_test.shape,x_valid.shape)
2 print(y_train.shape,y_test.shape,y_valid.shape)
```

```
(37, 3) (7, 3) (6, 3)
(37,) (7,) (6,)
```

7.Model Selection

```
1 def model_performance(model,model_name,x_train = x_train,y_train = y_train,x_test = x_test,y_test = y_
```

```

2
3     y_train_pred = model.predict(x_train)
4     y_test_pred = model.predict(x_test)
5     y_val_pred = model.predict(x_valid)
6
7     Training_Score = np.round(model.score(x_train,y_train),3)
8     Testing_Score = np.round(model.score(x_test,y_test),3)
9     Validation_score = np.round(model.score(x_valid,y_valid))
10
11     mse_training = np.round(mean_squared_error(y_train,y_train_pred),3)
12     mse_testing = np.round(mean_squared_error(y_test,y_test_pred),3)
13     mse_validation = np.round(mean_squared_error(y_valid,y_val_pred),3)
14
15     mae_training = np.round(mean_absolute_error(y_train,y_train_pred),3)
16     mae_testing = np.round(mean_absolute_error(y_test,y_test_pred),3)
17     mae_valid = np.round(mean_absolute_error(y_valid,y_val_pred),3)
18
19     r2_training = np.round(r2_score(y_train,y_train_pred),3)
20     r2_testing = np.round(r2_score(y_test,y_test_pred),3)
21     r2_valid = np.round(r2_score(y_valid,y_val_pred),3)
22
23     print("Model Performance for:",model_name)
24     print("")
25
26     print("Training Score:",Training_Score)
27     print("Testing Score:",Testing_Score)
28     print("Validation Score",Validation_score)
29     print("")
30
31     print("Training Data Mean Squared Error:",mse_training)
32     print("Testing Data Mean Squared Error:",mse_testing)
33     print("Validation Data Mean Squared Error:",mse_validation)
34
35     print("")
36
37     print("Training Data Mean Absolute Error:",mae_training)
38     print("Testing Data Mean Absolute Error:",mae_testing)
39     print("Validation Data Mean Absolute Error:",mae_valid)
40     print("")
41
42     print("Training Data r2_score:",r2_training)
43     print("Testing Data r2_score:",r2_testing)
44     print("Validation Data r2_score:",r2_valid)
45     print("")
46
47     print("Residual Analysis:")
48     plt.figure(figsize = (20,5))
49     plt.scatter(y_train,(y_train-y_train_pred),color = "red",label = 'Training Predictions')
50     plt.scatter(y_test,(y_test-y_test_pred),color = "green",label = 'Testing Predictions')
51     plt.scatter(y_valid,(y_valid-y_val_pred),color = 'blue',label = "Validation Predictions")
52     plt.legend()
53     plt.show()
54
55     return Training_Score,Testing_Score,Validation_score,mse_training,mse_testing,mse_validation,mae_t

```

▼ 1. Linear Regression

```

1  model1 = LinearRegression()
2  model1.fit(x_train,y_train)

```

```

LinearRegression()

```

```

1  lr_perf = model_performance(model1,model_name = model1)

```

Model Performance for: LinearRegression()


Training Score: 0.954
 Testing Score: 0.849
 Validation Score 1.0

Training Data Mean Squared Error: 0.05
 Testing Data Mean Squared Error: 0.056
 Validation Data Mean Squared Error: 0.053

Training Data Mean Absolute Error: 0.164
 Testing Data Mean Absolute Error: 0.189
 Validation Data Mean Absolute Error: 0.183

Training Data r2_score: 0.954
 Testing Data r2_score: 0.849
 Validation Data r2_score: 0.948

Residual Analysis:



2. Ridge

```
1 model2 = Ridge(alpha = 0.01)
2 model2.fit(x_train,y_train)
```

▼ Ridge
 Ridge(alpha=0.01)

```
1 ridge_perf = model_performance(model2,model2)
```

Model Performance for: Ridge(alpha=0.01)

Training Score: 0.954

3. KNeighborsRegressor

Testing Data Mean Squared Error: 0.056

```
1 model3 = KNeighborsRegressor(n_neighbors = 6)
2 model3.fit(x_train,y_train)
```

```
KNeighborsRegressor
KNeighborsRegressor(n_neighbors=6)
```

Validation Data r2 score: 0.948

```
1 knn_perf = model_performance(model3,model3)
```

Model Performance for: KNeighborsRegressor(n_neighbors=6)

Training Score: 0.898

Testing Score: 0.862

Validation Score 1.0

Training Data Mean Squared Error: 0.112

Testing Data Mean Squared Error: 0.051

Validation Data Mean Squared Error: 0.269

Training Data Mean Absolute Error: 0.252

Testing Data Mean Absolute Error: 0.202


Validation Data Mean Absolute Error: 0.437

Training Data r2_score: 0.898

Testing Data r2_score: 0.862

Validation Data r2_score: 0.735

Residual Analysis:



Linear Regression is giving best result

```
1 prediction = model1.predict(x_test)
```

```
1 prediction
```

```
array([-0.66439639, -0.46498393,  0.41686469, -0.28953794,  0.41544112,
       -1.65030452, -0.08646584])
```

```
1 # Get the mean and standard deviation from the scaler
```

```
2 mean_profit = scaler.mean_[0]
```

```
3 std_dev_profit = scaler.scale_[0]
```

```
4
```

```
5 scaled_prediction = prediction
```

```
6
```

```
7 # Reverse the scaling to get the prediction in the original units
```

```
8 original_prediction = (scaled_prediction * std_dev_profit) + mean_profit
```

```
9
```

```
10 print("Original prediction in dollars:", original_prediction)
```

```
Original prediction in dollars: [ 85502.50398527  93459.27699416 128645.99157053 100459.7619706
 128589.18988353  46163.70173114 108562.55837568]
```