

IT-641 Deep Learning

Lab 2

▼ House Price Prediction: Predict house prices using advanced regression techniques

SalePrice - the property's sale price in dollars. This is the target variable that you're trying to predict.

MSSubClass: The building class

MSZoning: The general zoning classification

LotFrontage: Linear feet of street connected to property

LotArea: Lot size in square feet

Street: Type of road access

Alley: Type of alley access

LotShape: General shape of property

LandContour: Flatness of the property

Utilities: Type of utilities available

LotConfig: Lot configuration

LandSlope: Slope of property

Neighborhood: Physical locations within Ames city limits

Condition1: Proximity to main road or railroad

Condition2: Proximity to main road or railroad (if a second is present)

BldgType: Type of dwelling

HouseStyle: Style of dwelling

OverallQual: Overall material and finish quality

OverallCond: Overall condition rating

YearBuilt: Original construction date

YearRemodAdd: Remodel date

RoofStyle: Type of roof

RoofMatl: Roof material

Exterior1st: Exterior covering on house

Exterior2nd: Exterior covering on house (if more than one material)

MasVnrType: Masonry veneer type

MasVnrArea: Masonry veneer area in square feet

ExterQual: Exterior material quality

ExterCond: Present condition of the material on the exterior

Foundation: Type of foundation

BsmtQual: Height of the basement

BsmtCond: General condition of the basement

BsmtExposure: Walkout or garden level basement walls

BsmtFinType1: Quality of basement finished area

BsmtFinSF1: Type 1 finished square feet

BsmtFinType2: Quality of second finished area (if present)

BsmtFinSF2: Type 2 finished square feet

BsmtUnfSF: Unfinished square feet of basement area

TotalBsmtSF: Total square feet of basement area

Heating: Type of heating

HeatingQC: Heating quality and condition

CentralAir: Central air conditioning

Electrical: Electrical system

1stFlrSF: First Floor square feet

2ndFlrSF: Second floor square feet

LowQualFinSF: Low quality finished square feet (all floors) GrLivArea: Above grade (ground) living area square feet

BsmtFullBath: Basement full bathrooms

BsmtHalfBath: Basement half bathrooms

FullBath: Full bathrooms above grade

HalfBath: Half baths above grade

Bedroom: Number of bedrooms above basement level

Kitchen: Number of kitchens

KitchenQual: Kitchen quality

TotRmsAbvGrd: Total rooms above grade (does not include bathrooms)

Functional: Home functionality rating

Fireplaces: Number of fireplaces

FireplaceQu: Fireplace quality

GarageType: Garage location

GarageYrBlt: Year garage was built

GarageFinish: Interior finish of the garage

GarageCars: Size of garage in car capacity

GarageArea: Size of garage in square feet

GarageQual: Garage quality

GarageCond: Garage condition

PavedDrive: Paved driveway

WoodDeckSF: Wood deck area in square feet

OpenPorchSF: Open porch area in square feet

EnclosedPorch: Enclosed porch area in square feet

3SsnPorch: Three season porch area in square feet

ScreenPorch: Screen porch area in square feet

PoolArea: Pool area in square feet

PoolQC: Pool quality

Fence: Fence quality

MiscFeature: Miscellaneous feature not covered in other categories

MiscVal: \$Value of miscellaneous feature

MoSold: Month Sold

YrSold: Year Sold

SaleType: Type of sale

SaleCondition: Condition of sale

It seems quite a lot of information in columns.

Tasks:

For the given datasets perform the following tasks:

1. Load Data and Find out if it has any missing values
2. Correct the missing values
3. Identify and encode necessary features
4. Identify and normalize necessary features

5. Split the dataset into train set (75%) test set (15%) an validation set (10%)
6. For Regression Task: Write a code manually for Linear Regression and compare the results with sklearn's linear regression model.
7. For Classification Task: Write manual code for logistic regression using Gradient Descent and compare it with sklearn's Logistic Regression
8. Use all the respective model performance criteria and compare to model
9. Discuss under-fitting and overfitting based upon results

1. Loading Required Libraries

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sb
5 from sklearn.preprocessing import LabelEncoder, StandardScaler
6 from sklearn.model_selection import train_test_split, GridSearchCV
7 import warnings
8 warnings.filterwarnings(action = 'ignore')
9 from sklearn.preprocessing import LabelEncoder, StandardScaler
10 from sklearn.linear_model import LinearRegression
11 from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score, cross_val_predict
12 from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
13
```

2. Load Dataset

```
1 data = pd.read_csv("https://raw.githubusercontent.com/Jatansahu/DEEP_LEARNING_ASSIGNMENTS/main/LAB_02/
```

```
1 # Previewing data
2 data.head()
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	Mis
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	

5 rows × 81 columns

Looking at the above dataset our target variable is the column "SalePrice"

3. Looking for Null values

```
1 data.info()
```

```


40 heatingQC      1460 non-null object
41 CentralAir     1460 non-null object
42 Electrical     1459 non-null object
43 1stFlrSF       1460 non-null int64
44 2ndFlrSF       1460 non-null int64
45 LowQualFinSF   1460 non-null int64
46 GrLivArea      1460 non-null int64
47 BsmtFullBath   1460 non-null int64
48 BsmtHalfBath   1460 non-null int64
49 FullBath       1460 non-null int64
50 HalfBath       1460 non-null int64
51 BedroomAbvGr  1460 non-null int64
52 KitchenAbvGr  1460 non-null int64
53 KitchenQual    1460 non-null object
54 TotRmsAbvGrd  1460 non-null int64
55 Functional     1460 non-null object
56 Fireplaces     1460 non-null int64
57 FireplaceQu    770 non-null object
58 GarageType     1379 non-null object
59 GarageYrBlt    1379 non-null float64
60 GarageFinish   1379 non-null object
61 GarageCars     1460 non-null int64
62 GarageArea     1460 non-null int64
63 GarageQual     1379 non-null object
64 GarageCond     1379 non-null object
65 PavedDrive     1460 non-null object
66 WoodDeckSF     1460 non-null int64
67 OpenPorchSF    1460 non-null int64
68 EnclosedPorch  1460 non-null int64
69 3SsnPorch      1460 non-null int64
70 ScreenPorch    1460 non-null int64
71 PoolArea       1460 non-null int64
72 PoolQC         7 non-null object
73 Fence          281 non-null object
74 MiscFeature    54 non-null object
75 MiscVal        1460 non-null int64
76 MoSold         1460 non-null int64
77 YrSold         1460 non-null int64
78 SaleType       1460 non-null object
79 SaleCondition  1460 non-null object
80 SalePrice      1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```

1 import missingno as msno
2 msno.bar(data)

```



Dropping those columns which have less than 25% values



```

1 # Calculate the threshold for null values
2 threshold = len(data) * 0.25 # 25% of total rows
3
4 # Iterate through the columns and drop those with null percentages greater than the threshold
5 for column in data.columns:
6     if data[column].isnull().sum() > threshold:
7         data.drop(column, axis=1, inplace=True)
8 print(msno.bar(data))

```



Double-click (or enter) to edit



We still have some missing values in our dataset we will look into further



4.Basic EDA



1.Descriptive Analysis



```
1 data.describe().style.background_gradient()
```

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFi
count	1460.000000	1460.000000	1201.000000	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	1452.000000	1460.00
mean	730.500000	56.897260	70.049958	10516.828082	6.099315	5.575342	1971.267808	1984.865753	103.685262	443.63
std	421.610009	42.300571	24.284752	9981.264932	1.382997	1.112799	30.202904	20.645407	181.066207	456.09
min	1.000000	20.000000	21.000000	1300.000000	1.000000	1.000000	1872.000000	1950.000000	0.000000	0.00
25%	365.750000	20.000000	59.000000	7553.500000	5.000000	5.000000	1954.000000	1967.000000	0.000000	0.00
50%	730.500000	50.000000	69.000000	9478.500000	6.000000	5.000000	1973.000000	1994.000000	0.000000	383.50
75%	1095.250000	70.000000	80.000000	11601.500000	7.000000	6.000000	2000.000000	2004.000000	166.000000	712.25
max	1460.000000	190.000000	313.000000	215245.000000	10.000000	9.000000	2010.000000	2010.000000	1600.000000	5644.00

2 Correlation

```
1 correlation_matrix = data.corr()
2 print("Correlation Matrix:")
3 (correlation_matrix.style.background_gradient())
```

Correlation Matrix:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea	BsmtFinS
Id	1.000000	0.011156	-0.010601	-0.033226	-0.028365	0.012609	-0.012713	-0.021998	-0.050298	-0.0050
MSSubClass	0.011156	1.000000	-0.386347	-0.139781	0.032628	-0.059316	0.027850	0.040581	0.022936	-0.0698
LotFrontage	-0.010601	-0.386347	1.000000	0.426095	0.251646	-0.059213	0.123349	0.088866	0.193458	0.2336
LotArea	-0.033226	-0.139781	0.426095	1.000000	0.105806	-0.005636	0.014228	0.013788	0.104160	0.2141
OverallQual	-0.028365	0.032628	0.251646	0.105806	1.000000	-0.091932	0.572323	0.550684	0.411876	0.2396
OverallCond	0.012609	-0.059316	-0.059213	-0.005636	-0.091932	1.000000	-0.375983	0.073741	-0.128101	-0.0462
YearBuilt	-0.012713	0.027850	0.123349	0.014228	0.572323	-0.375983	1.000000	0.592855	0.315707	0.2495
YearRemodAdd	-0.021998	0.040581	0.088866	0.013788	0.550684	0.073741	0.592855	1.000000	0.179618	0.1284
MasVnrArea	-0.050298	0.022936	0.193458	0.104160	0.411876	-0.128101	0.315707	0.179618	1.000000	0.2647
BsmtFinSF1	-0.005024	-0.069836	0.233633	0.214103	0.239666	-0.046231	0.249503	0.128451	0.264736	1.0000
BsmtFinSF2	-0.005968	-0.065649	0.049900	0.111170	-0.059119	0.040229	-0.049107	-0.067759	-0.072319	-0.0501
BsmtUnfSF	-0.007940	-0.140759	0.132644	-0.002618	0.308159	-0.136841	0.149040	0.181133	0.114442	-0.4952
TotalBsmtSF	-0.015415	-0.238518	0.392075	0.260833	0.537808	-0.171098	0.391452	0.291066	0.363936	0.5223
1stFlrSF	0.010496	-0.251758	0.457181	0.299475	0.476224	-0.144203	0.281986	0.240379	0.344501	0.4458
2ndFlrSF	0.005590	0.307886	0.080177	0.050986	0.295493	0.028942	0.010308	0.140024	0.174561	-0.1370
LowQualFinSF	-0.044230	0.046474	0.038469	0.004779	-0.030429	0.025494	-0.183784	-0.062419	-0.069071	-0.0645
GrLivArea	0.008273	0.074853	0.402797	0.263116	0.593007	-0.079686	0.199010	0.287389	0.390857	0.2081
BsmtFullBath	0.002289	0.003491	0.100949	0.158155	0.111098	-0.054942	0.187599	0.119470	0.085310	0.6492
BsmtHalfBath	-0.020155	-0.002333	-0.007234	0.048046	-0.040150	0.117821	-0.038162	-0.012337	0.026673	0.0674
FullBath	0.005587	0.131608	0.198769	0.126031	0.550600	-0.194149	0.468271	0.439046	0.276833	0.0585
HalfBath	0.006784	0.177354	0.053532	0.014259	0.273458	-0.060769	0.242656	0.183331	0.201444	0.0042
BedroomAbvGr	0.037719	-0.023438	0.263170	0.119690	0.101676	0.012980	-0.070651	-0.040581	0.102821	-0.1073
KitchenAbvGr	0.002951	0.281721	-0.006069	-0.017784	-0.183882	-0.087001	-0.174800	-0.149598	-0.037610	-0.0810
TotRmsAbvGrd	0.027239	0.040380	0.352096	0.190015	0.427452	-0.057583	0.095589	0.191740	0.280682	0.0443
Fireplaces	-0.019772	-0.045569	0.266639	0.271364	0.396765	-0.023820	0.147716	0.112581	0.249070	0.2600
GarageYrBlt	0.000072	0.085072	0.070250	-0.024947	0.547766	-0.324297	0.825667	0.642277	0.252691	0.1534
GarageCars	0.016570	-0.040110	0.285691	0.154871	0.600671	-0.185758	0.537850	0.420622	0.364204	0.2240
GarageArea	0.017634	-0.098672	0.344997	0.180403	0.562022	-0.151521	0.478954	0.371600	0.373066	0.2969
WoodDeckSF	-0.029643	-0.012579	0.088521	0.171698	0.238923	-0.003334	0.224880	0.205726	0.159718	0.2043

3.Feature selection (Numerical)

```

1 # Set the correlation threshold for feature selection
2 correlation_threshold = 0.3
3
4 print("Before Feature Selection ",data.shape)
5
6 # Select features with correlation above the threshold
7 selected_features = []
8 for column in correlation_matrix.columns:
9     if abs(correlation_matrix['SalePrice'][column]) < correlation_threshold:
10         data = data.drop(columns=column)
11     else:
12         selected_features.append(column)
13 print("Selected Features",selected_features)
14 print("After Feature Selection ",data.shape)


```

Before Feature Selection (1460, 76)
Selected Features ['LotFrontage', 'OverallQual', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1', 'TotalBsmtSF', '1stFlrSF']
After Feature Selection (1460, 57)

4. Checking Skewness of Target Variable

```
1 import seaborn as sns
2 sns.distplot(data['SalePrice'], color='g', bins=100, hist_kws={'alpha' : 0.4})
```

<Axes: xlabel='SalePrice', ylabel='Density'>



Salesprice is right skewed means mean > median

5. Preprocessing

```
1 data.head()
```

	MSZoning	LotFrontage	Street	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	...	GarageCars
0	RL	65.0	Pave	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	...	2
1	RL	80.0	Pave	Reg	Lvl	AllPub	FR2	Gtl	Veenker	Feedr	...	2
2	RL	68.0	Pave	IR1	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	...	2
3	RL	60.0	Pave	IR1	Lvl	AllPub	Corner	Gtl	Crawfor	Norm	...	3
4	RL	84.0	Pave	IR1	Lvl	AllPub	FR2	Gtl	NoRidge	Norm	...	3

5 rows × 57 columns

1. Dropping Unnecessary columns

Dropping Id column which is not useful. We already deleted in correlation part

Scaling feature

```
1 SalePrice_data = data[["SalePrice"]] # Extracting the "SalePrice" column as a separate DataFrame
2 scaler = StandardScaler()
3 scaled_SalePrice = scaler.fit_transform(SalePrice_data.values.reshape(-1,1))
```

```
1 # Dropping profit column from dataset
2 data.drop(["SalePrice"],1,inplace = True)
3 data['scaled_SalePrice'] = scaled_SalePrice
```

```
1 # # Get the mean and standard deviation from the scaler
2 # mean_profit = scaler.mean_[0]
3 # std_dev_profit = scaler.scale_[0]
4
5 # scaled_prediction = 2.01120333
```



```
6
7 # # Reverse the scaling to get the prediction in the original units
8 # original_prediction = (scaled_prediction * std_dev_profit) + mean_profit
9
10 # print("Original prediction in dollars:", original_prediction)

1 sc = StandardScaler()
2 for col in data.columns:
3     if data[col].dtype in ['int64' , 'float64']:
4         data[col] = sc.fit_transform(data[col].values.reshape(-1,1))
```

2.Converting Categorical Variables into their corresponding form

▼ A).Cardinality

"Cardinality" means the number of unique values in a column

```
1 for col in data.columns:
2     if data[col].dtype=='object':
3         print()
4         print(col)
5         print('Number of unique values :',data[col].nunique())
6         print('Sample unique values :',data[col].unique()[:5])
```

```
SaleCondition
Number of unique values : 6
Sample unique values : ['Normal' 'Abnorml' 'Partial' 'AdjLand' 'Alloca']
```

▼ B) Encoding all the categorical variable using labelencoder

```
1 categorical_cols = [cname for cname in data.columns
2                     if data[cname].dtype == "object"]
3 # categorical_cols

1 #encoding Embarked column
2 le = LabelEncoder()
3 for col in data.columns:
4     if data[col].dtype=='object':
5         data[col] = le.fit_transform(data[col])

1 sc = StandardScaler()
2 for col in data.columns:
3     if data[col].dtype in ['int64' , 'float64']:
4         data[col] = sc.fit_transform(data[col].values.reshape(-1,1))
```

```
1 data.info()

1  LotFrontage      1201 non-null  float64
2  Street           1460 non-null  float64
3  LotShape         1460 non-null  float64
4  LandContour      1460 non-null  float64
5  Utilities        1460 non-null  float64
6  LotConfig        1460 non-null  float64
7  LandSlope        1460 non-null  float64
8  Neighborhood     1460 non-null  float64
9  Condition1       1460 non-null  float64
10 Condition2       1460 non-null  float64
11 BldgType         1460 non-null  float64
12 HouseStyle       1460 non-null  float64
13 OverallQual      1460 non-null  float64
14 YearBuilt        1460 non-null  float64
15 YearRemodAdd     1460 non-null  float64
16 RoofStyle        1460 non-null  float64
17 RoofMatl         1460 non-null  float64
18 Exterior1st      1460 non-null  float64
19 Exterior2nd      1460 non-null  float64
20 MasVnrType       1460 non-null  float64
21 MasVnrArea       1452 non-null  float64
22 ExterQual        1460 non-null  float64
23 ExterCond        1460 non-null  float64
24 Foundation       1460 non-null  float64
25 BsmtQual         1460 non-null  float64
26 BsmtCond         1460 non-null  float64
27 BsmtExposure     1460 non-null  float64
28 BsmtFinType1     1460 non-null  float64
29 BsmtFinSF1       1460 non-null  float64
30 BsmtFinType2     1460 non-null  float64
31 TotalBsmtSF      1460 non-null  float64
32 Heating          1460 non-null  float64
33 HeatingQC        1460 non-null  float64
34 CentralAir       1460 non-null  float64
35 Electrical       1460 non-null  float64
36 1stFlrSF         1460 non-null  float64
37 2ndFlrSF         1460 non-null  float64
38 GrLivArea        1460 non-null  float64
39 FullBath         1460 non-null  float64
40 KitchenQual      1460 non-null  float64
41 TotRmsAbvGrd     1460 non-null  float64
42 Functional       1460 non-null  float64
43 Fireplaces       1460 non-null  float64
44 GarageType       1460 non-null  float64
45 GarageYrBlt      1379 non-null  float64
46 GarageFinish     1460 non-null  float64
47 GarageCars       1460 non-null  float64
48 GarageArea       1460 non-null  float64
49 GarageQual       1460 non-null  float64
50 GarageCond       1460 non-null  float64
```

▼ Data imputation

```
1 def auto_data_impute(data,get_rid_percent=2):
2     for x,y in data.isnull().sum().items():
3         percent = y/data.shape[0]
4         if percent <= get_rid_percent/100:
5             data[x] = data[x].fillna(data[x].mean())
6         else:
7             print("removed column : ",x)
8             data = data.drop([x],axis=1)
9     print("Data imputation successfull")
10    return data
```

```
1 data = auto_data_impute(data,get_rid_percent=2)
```

removed column : LotFrontage
removed column : GarageYrBlt
Data imputation successfull

▼ 6.Splitting the dataset

```
1 data
```

	MSZoning	Street	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Condition2	...	GarageC
0	-0.045532	0.064238	0.750731	0.314667	-0.02618	0.604670	-0.225716	-1.206215	-0.036289	-0.03174	...	0.311
1	-0.045532	0.064238	0.750731	0.314667	-0.02618	-0.628316	-0.225716	1.954302	-1.188074	-0.03174	...	0.311
2	-0.045532	0.064238	-1.378933	0.314667	-0.02618	0.604670	-0.225716	-1.206215	-0.036289	-0.03174	...	0.311
3	-0.045532	0.064238	-1.378933	0.314667	-0.02618	-1.861302	-0.225716	-1.039872	-0.036289	-0.03174	...	1.650
4	-0.045532	0.064238	-1.378933	0.314667	-0.02618	-0.628316	-0.225716	0.457215	-0.036289	-0.03174	...	1.650
...
1455	-0.045532	0.064238	0.750731	0.314667	-0.02618	0.604670	-0.225716	-0.707186	-0.036289	-0.03174	...	0.311
1456	-0.045532	0.064238	0.750731	0.314667	-0.02618	0.604670	-0.225716	0.290872	-0.036289	-0.03174	...	0.311
1457	-0.045532	0.064238	0.750731	0.314667	-0.02618	0.604670	-0.225716	-1.039872	-0.036289	-0.03174	...	-1.026
1458	-0.045532	0.064238	0.750731	0.314667	-0.02618	0.604670	-0.225716	-0.041814	-0.036289	-0.03174	...	-1.026
1459	-0.045532	0.064238	0.750731	0.314667	-0.02618	0.604670	-0.225716	-0.873529	-0.036289	-0.03174	...	-1.026

1460 rows × 55 columns



```
1 x = data.iloc[:,54]
2 y = data['scaled_SalePrice']
```

```
1 y
```

0 0.347273
1 0.007288
2 0.536154
3 -0.515281
4 0.869843
...
1455 -0.074560
1456 0.366161
1457 1.077611
1458 -0.488523
1459 -0.420841
Name: scaled_SalePrice, Length: 1460, dtype: float64

```
1 x.head()
```

Neighborhood	Condition1	Condition2	...	GarageFinish	GarageCars	GarageArea	GarageQual	GarageCond	PavedDrive	WoodDeckSF	OpenPorchSF
-1.206215	-0.036289	-0.03174	...	-0.318475	0.311725	0.351000	0.11211	0.0689	0.289745	-0.752176	0
1.954302	-1.188074	-0.03174	...	-0.318475	0.311725	-0.060731	0.11211	0.0689	0.289745	1.626195	-0
-1.206215	-0.036289	-0.03174	...	-0.318475	0.311725	0.631726	0.11211	0.0689	0.289745	-0.752176	-0
-1.039872	-0.036289	-0.03174	...	0.801942	1.650307	0.790804	0.11211	0.0689	0.289745	-0.752176	-0
0.457215	-0.036289	-0.03174	...	-0.318475	1.650307	1.698485	0.11211	0.0689	0.289745	0.780197	0

Training set - 75%

Testing set - 15%

Validation set - 10%

```
1 x_train,x_part,y_train,y_part = train_test_split(x,y,test_size = 0.25,random_state = 1)
2 x_test,x_valid,y_test,y_valid = train_test_split(x_part,y_part,test_size = 0.4,random_state = 1)

1 print(x_train.shape,x_test.shape,x_valid.shape)
2 print(y_train.shape,y_test.shape,y_valid.shape)

(1095, 54) (219, 54) (146, 54)
(1095,) (219,) (146,)
```

7.Model Selection

▾ A) Linear regression from scratch

```
1 class LinearRegression:
2     def __init__(self, learning_rate=0.01, num_iterations=1000):
3         self.learning_rate = learning_rate
4         self.num_iterations = num_iterations
5         self.weights = None
6         self.bias = None
7
8     def fit(self, X, y):
9         num_samples, num_features = X.shape
10        self.weights = np.zeros(num_features)
11        self.bias = 0
12
13        # Gradient descent optimization
14        for _ in range(self.num_iterations):
15            predicted = np.dot(X, self.weights) + self.bias
16            error = y - predicted
17
18            # Update weights and bias using gradients
19            self.weights += (self.learning_rate / num_samples) * np.dot(X.T, error)
20            self.bias += (self.learning_rate / num_samples) * np.sum(error)
21
22        def predict(self, X):
23            return np.dot(X, self.weights) + self.bias
24
25 # Create a linear regression model
26 model = LinearRegression(learning_rate=0.01, num_iterations=1000)
27
28 # Fit the model to the data
29 model.fit(x_train, y_train)

1 # Calculate Mean Squared Error (MSE)
2 def mean_squared_error(y_test,y_pred ):
3     return np.mean((y_test - y_pred) ** 2)
4
5 # Calculate Root Mean Squared Error (RMSE)
6 def root_mean_squared_error(y_test, y_pred):
7     return np.sqrt(mean_squared_error(y_test, y_pred))
8
```

```

8
9 # Calculate R-squared (coefficient of determination)
10 def r_squared(y_test, y_pred):
11     y_mean = np.mean(y_test)
12     ss_total = np.sum((y_test - y_mean) ** 2)
13     ss_residual = np.sum((y_test - y_pred) ** 2)
14     return 1 - (ss_residual / ss_total)
15

```

```

1 # For testing
2 # Make predictions
3 y_pred = model.predict(x_test)

```

```

1 # Calculate metrics
2 mse = mean_squared_error(y_test, y_pred)
3 rmse = root_mean_squared_error(y_test, y_pred)
4 r2 = r_squared(y_test, y_pred)
5
6 print("Mean Squared Error:", mse)
7 print("Root Mean Squared Error:", rmse)
8 print("R-squared:", r2)

```

```

Mean Squared Error: 0.18648684238029734
Root Mean Squared Error: 0.43184122357678795
R-squared: 0.8466170598362222

```

▼ USING SCIKIT LEARN LIBRARY

```

1 def model_performance(model,model_name,x_train = x_train,y_train = y_train,x_test = x_test,y_test = y_
2
3     y_train_pred = model.predict(x_train)
4     y_test_pred = model.predict(x_test)
5     y_val_pred = model.predict(x_valid)
6
7     # Training_Score = np.round(model.score(x_train,y_train),3)
8     # Testing_Score = np.round(model.score(x_test,y_test),3)
9     # Validation_score = np.round(model.score(x_valid,y_valid))
10
11     # mse_training = np.round(mean_squared_error(y_train,y_train_pred),3)
12     mse_testing = np.round(mean_squared_error(y_test,y_test_pred),3)
13     # mse_validation = np.round(mean_squared_error(y_valid,y_val_pred),3)
14
15     # mae_training = np.round(mean_absolute_error(y_train,y_train_pred),3)
16     mae_testing = np.round(mean_absolute_error(y_test,y_test_pred),3)
17     # mae_valid = np.round(mean_absolute_error(y_valid,y_val_pred),3)
18
19     # r2_training = np.round(r2_score(y_train,y_train_pred),3)
20     r2_testing = np.round(r2_score(y_test,y_test_pred),3)
21     # r2_valid = np.round(r2_score(y_valid,y_val_pred),3)
22
23     print("Model Performance for:",model_name)
24     print("")
25
26     # print("Training Score:",Training_Score)
27     # print("Testing Score:",Testing_Score)
28     # print("Validation Score",Validation_score)
29     # print("")
30
31     # print("Training Data Mean Squared Error:",mse_training)
32     print("Testing Data Mean Squared Error:",mse_testing)
33     # print("Validation Data Mean Squared Error:",mse_validation)
34
35     print("")
36
37     # print("Training Data Mean Absolute Error:",mae_training)
38     print("Testing Data Mean Absolute Error:",mae_testing)
39     # print("Validation Data Mean Absolute Error:".mae valid)

```

```

39     # print("Training Data r2_score:",r2_training)
40     print("")
41
42     # print("Training Data r2_score:",r2_training)
43     print("Testing Data r2_score:",r2_testing)
44     # print("Validation Data r2_score:",r2_valid)
45     print("")
46
47     print("Residual Analysis:")
48     # plt.figure(figsize = (20,5))
49     # # plt.scatter(y_train,(y_train-y_train_pred),color = "red",label = 'Training Predictions')
50     # plt.scatter(y_test,(y_test-y_test_pred),color = "green",label = 'Testing Predictions')
51     # # plt.scatter(y_valid,(y_valid-y_val_pred),color = 'blue',label = "Validation Predictions")
52     # plt.legend()
53     # plt.show()
54
55     return mse_testing,mae_testing,r2_testing

1  model1 = LinearRegression()
2  model1.fit(x_train,y_train)

```

▼ Comparisation

```
1 lr_perf = model_performance(model1,model_name = model1)
```

Model Performance for: <__main__.LinearRegression object at 0x7c1df46b4850>

Testing Data Mean Squared Error: 0.186

Testing Data Mean Absolute Error: 0.248

Testing Data r2_score: 0.847

Residual Analysis:

Linear Regression written in scrach (For testing set)

```

Mean Squared Error: 0.2155905940198445
Root Mean Squared Error: 0.4643173419331271
R-squared: 0.822679612350427

```

1