

IT-641 Deep Learning

Lab 5

▸ INSTRUCTIONS:

↳ 4 cells hidden

DATASET01 - Monthly ocean dataset

▸ 1. Importing Required Libraries

[] ↳ 1 cell hidden

▸ 2. Data Understanding

[] ↳ 4 cells hidden

▸ Seprating Data Frames

[] ↳ 11 cells hidden

▸ TASK 1-3

↳ 1 cell hidden

▸ Lag Feature

[] ↳ 11 cells hidden

▸ FOR GCAG DATASET

[] ↳ 12 cells hidden

▾ Task5. Plot graphs showing the result for all the models.

```

1 # Predictions for all models
2 linear_reg_preds = linear_reg_model.predict(X_test)
3 poly_reg_preds = [poly_reg_model.predict(poly_features.transform(X_test)) for _, poly_reg_model in pol
4 ann_preds = ann_model.predict(X_test)
5
6 # Plotting the results
7 plt.figure(figsize=(12, 8))
8
9 # Actual vs. Predicted for Linear Regression
10 plt.subplot(2, 2, 1)
11 plt.scatter(X_test, y_test, color='blue', label='Actual')
12 plt.scatter(X_test, linear_reg_preds, color='red', label='Linear Regression')
13 plt.xlabel('X_test')
14 plt.ylabel('y_test')
15 plt.title('Actual vs. Predicted (Linear Regression)')
16 plt.legend()
17

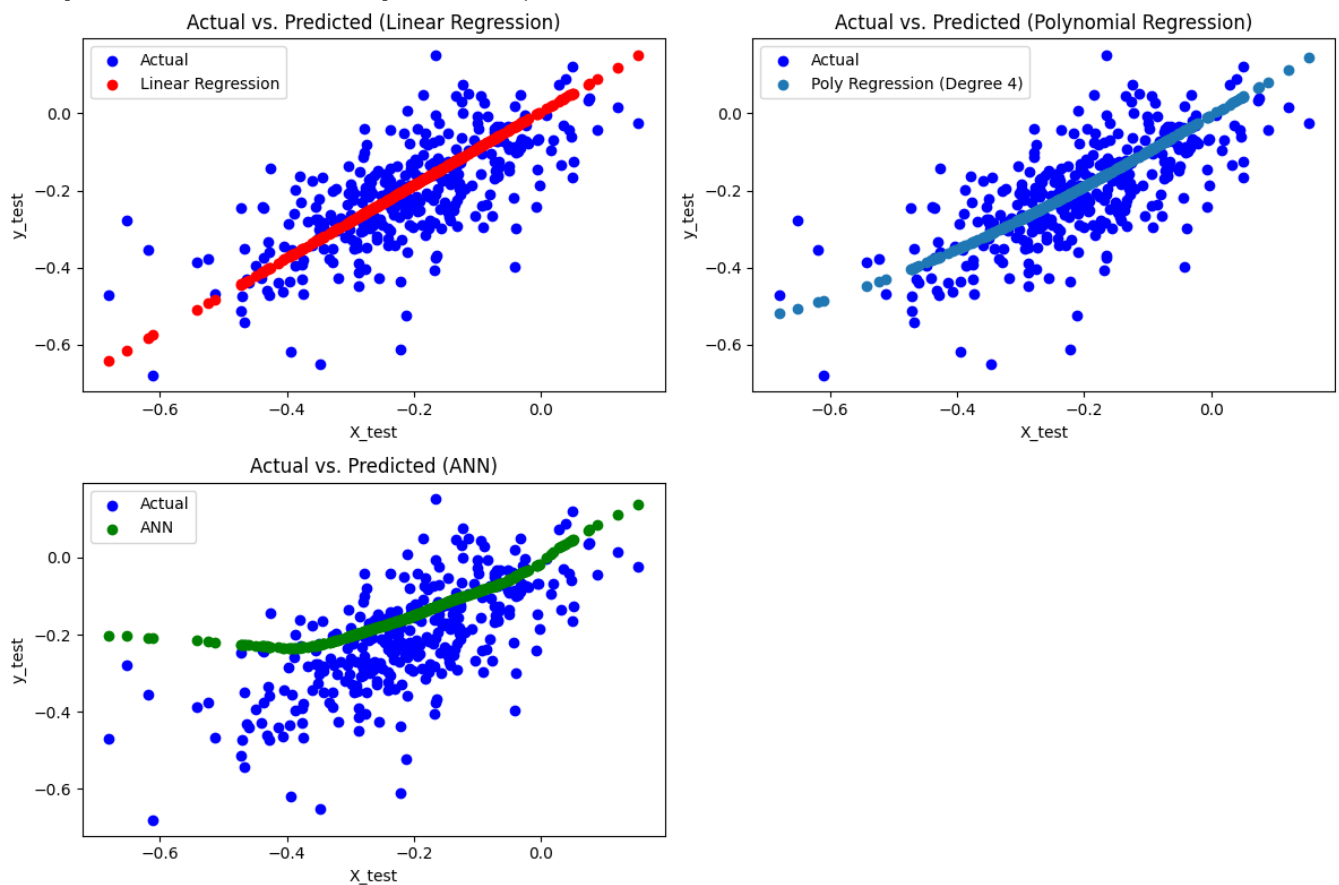
```

```

18 # Actual vs. Predicted for Polynomial Regression
19 plt.subplot(2, 2, 2)
20 plt.scatter(X_test, y_test, color='blue', label='Actual')
21 for degree, poly_reg_model in poly_reg_models:
22     plt.scatter(X_test, poly_reg_model.predict(poly_features.transform(X_test)), label=f'Poly Regression {degree}')
23 plt.xlabel('X_test')
24 plt.ylabel('y_test')
25 plt.title('Actual vs. Predicted (Polynomial Regression)')
26 plt.legend()
27
28 # Actual vs. Predicted for ANN
29 plt.subplot(2, 2, 3)
30 plt.scatter(X_test, y_test, color='blue', label='Actual')
31 plt.scatter(X_test, ann_preds, color='green', label='ANN')
32 plt.xlabel('X_test')
33 plt.ylabel('y_test')
34 plt.title('Actual vs. Predicted (ANN)')
35 plt.legend()
36
37 plt.tight_layout()
38 plt.show()
39
40

```

11/11 [=====] - 0s 2ms/step



▼ 6. Compare RMSE, MAE, MSE for the created models.

```

1 from sklearn.metrics import mean_squared_error, mean_absolute_error
2
3 # Function to calculate RMSE, MAE, and MSE
4 def calculate_metrics(y_true, y_pred):
5     rmse = np.sqrt(mean_squared_error(y_true, y_pred))

```

```

6     mae = mean_absolute_error(y_true, y_pred)
7     mse = mean_squared_error(y_true, y_pred)
8     return rmse, mae, mse
9
10 # Calculate metrics for Linear Regression
11 linear_reg_rmse, linear_reg_mae, linear_reg_mse = calculate_metrics(y_test, linear_reg_preds)
12
13 # Calculate metrics for Polynomial Regression models
14 poly_reg_metrics = []
15 for degree, poly_reg_model in poly_reg_models:
16     poly_reg_preds = poly_reg_model.predict(poly_features.transform(X_test))
17     rmse, mae, mse = calculate_metrics(y_test, poly_reg_preds)
18     poly_reg_metrics.append((degree, rmse, mae, mse))
19
20 # Calculate metrics for ANN
21 ann_rmse, ann_mae, ann_mse = calculate_metrics(y_test, ann_preds)
22
23 # Print and compare metrics
24 print(f"Linear Regression - RMSE: {linear_reg_rmse}, MAE: {linear_reg_mae}, MSE: {linear_reg_mse}")
25
26 for degree, rmse, mae, mse in poly_reg_metrics:
27     print(f"Polynomial Regression (Degree {degree}) - RMSE: {rmse}, MAE: {mae}, MSE: {mse}")
28
29 print(f"ANN - RMSE: {ann_rmse}, MAE: {ann_mae}, MSE: {ann_mse}")
30
Linear Regression - RMSE: 0.10685999521568819, MAE: 0.08251771071017241, MSE: 0.011419058577496902
Polynomial Regression (Degree 4) - RMSE: 0.10337531742011694, MAE: 0.08009041191718375, MSE: 0.010686456251709931
ANN - RMSE: 0.12506471580165127, MAE: 0.0958508650961165, MSE: 0.0156411831385478

```

▼ 7 Compare the results and justify which one is better.

Polynomial regression with degree= 4 will works better in this as it gives us less error in other cases. Data is non linear in this case.

1

▼ FOR GISTEMP DATASET

```

1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(X_gistemp, y_gistemp,
3                                                     test_size=0.2, shuffle=False)
4 X_train = np.array(X_train).reshape(-1, 1)
5 y_train = list(y_train)
6 X_test = np.array(X_test).reshape(-1, 1)
7 y_test = list(y_test)

```

▼ TASK 4 Model train

4. Create and train LinearRegression, PolynomialRegression (with three different number of polynomial features) and ANN model for both datasets.

```

1 # Reshape data to match model input requirements
2 X_train = np.array(X_train).reshape(-1, 1)
3 y_train = np.array(y_train) # Convert y_train to a NumPy array
4 X_test = np.array(X_test).reshape(-1, 1)
5 y_test = np.array(y_test) # Convert y_test to a NumPy array

```

```

1 #Checking which degree is working best
2 poly = PolynomialFeatures(2)
3 new_x = poly.fit_transform(X_train)
4
5 model = LinearRegression()

```

```
6 model.fit(new_x, y_train)
7
8 new_x = poly.fit_transform(X_test)
9
10 y_pred = pd.Series(model.predict(new_x))
11 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
12 rmse
0.12426652001566257
```

```
1 #Checking which degree is working best
2 poly = PolynomialFeatures(3)
3 new_x = poly.fit_transform(X_train)
4
5 model = LinearRegression()
6 model.fit(new_x, y_train)
7
8 new_x = poly.fit_transform(X_test)
9
10 y_pred = pd.Series(model.predict(new_x))
11 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
12 rmse
0.12149788572830693
```

```
1 #Checking which degree is working best
2 poly = PolynomialFeatures(4)
3 new_x = poly.fit_transform(X_train)
4
5 model = LinearRegression()
6 model.fit(new_x, y_train)
7
8 new_x = poly.fit_transform(X_test)
9
10 y_pred = pd.Series(model.predict(new_x))
11 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
12 rmse
0.1209652282880449
```

```
1 #Checking which degree is working best
2 poly = PolynomialFeatures(5)
3 new_x = poly.fit_transform(X_train)
4
5 model = LinearRegression()
6 model.fit(new_x, y_train)
7
8 new_x = poly.fit_transform(X_test)
9
10 y_pred = pd.Series(model.predict(new_x))
11 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
12 rmse
0.12095827506545165
```

```
1 #Checking which degree is working best
2 poly = PolynomialFeatures(6)
3 new_x = poly.fit_transform(X_train)
4
5 model = LinearRegression()
6 model.fit(new_x, y_train)
7
8 new_x = poly.fit_transform(X_test)
9
10 y_pred = pd.Series(model.predict(new_x))
11 rmse = np.sqrt(mean_squared_error(y_test, y_pred))
12 rmse
0.12162328362286128
```

In this dataset degree 5 works well we will do further for degree 5

```
1 # Create and train Linear Regression model
2 linear_reg_model = LinearRegression()
3 linear_reg_model.fit(X_train, y_train)
4
5 # Create and train Polynomial Regression models with different degrees
6 poly_degrees = [5]
7 poly_reg_models = []
8
9
10 for degree in poly_degrees:
11     poly_features = PolynomialFeatures(degree=degree)
12     X_train_poly = poly_features.fit_transform(X_train)
13     X_test_poly = poly_features.transform(X_test)
14
15     poly_reg_model = LinearRegression()
16     poly_reg_model.fit(X_train_poly, y_train)
17
18     poly_reg_models.append((degree, poly_reg_model))
19
20 # Create and train an Artificial Neural Network (ANN) model
21 ann_model = keras.Sequential([
22     layers.Dense(64, activation='relu', input_shape=(1,)),
23     layers.Dense(32, activation='relu'),
24     layers.Dense(1) # Output layer
25 ])
26
27 # Compile the ANN model
28 ann_model.compile(optimizer='adam', loss='mean_squared_error')
29
30 # Train the ANN model
31 ann_model.fit(X_train, y_train, epochs=100, batch_size=32, validation_split=0.2)
```

```

Epoch 94/100
33/33 [=====] - 0s 3ms/step - loss: 0.0140 - val_loss: 0.0228
Epoch 95/100
33/33 [=====] - 0s 3ms/step - loss: 0.0140 - val_loss: 0.0290
Epoch 96/100
33/33 [=====] - 0s 3ms/step - loss: 0.0141 - val_loss: 0.0223
Epoch 97/100
33/33 [=====] - 0s 3ms/step - loss: 0.0139 - val_loss: 0.0274
Epoch 98/100
33/33 [=====] - 0s 4ms/step - loss: 0.0140 - val_loss: 0.0238
Epoch 99/100
33/33 [=====] - 0s 3ms/step - loss: 0.0140 - val_loss: 0.0285
Epoch 100/100
33/33 [=====] - 0s 3ms/step - loss: 0.0141 - val_loss: 0.0237
<keras.src.callbacks.History at 0x7a3579a81f30>

```


Task5. Plot graphs showing the result for all the models.

```

1 import matplotlib.pyplot as plt
2
3 # Predictions for all models
4 linear_reg_preds = linear_reg_model.predict(X_test)
5 poly_reg_preds = [poly_reg_model.predict(poly_features.transform(X_test)) for _, poly_reg_model in pol
6 ann_preds = ann_model.predict(X_test)
7
8 # Plotting the results
9 plt.figure(figsize=(12, 8))
10
11 # Actual vs. Predicted for Linear Regression
12 plt.subplot(2, 2, 1)
13 plt.scatter(X_test, y_test, color='blue', label='Actual')
14 plt.scatter(X_test, linear_reg_preds, color='red', label='Linear Regression')
15 plt.xlabel('X_test')
16 plt.ylabel('y_test')
17 plt.title('Actual vs. Predicted (Linear Regression)')
18 plt.legend()
19
20 # Actual vs. Predicted for Polynomial Regression
21 plt.subplot(2, 2, 2)
22 plt.scatter(X_test, y_test, color='blue', label='Actual')
23 for degree, poly_reg_model in poly_reg_models:
24     plt.scatter(X_test, poly_reg_model.predict(poly_features.transform(X_test)), label=f'Poly Regressi
25 plt.xlabel('X_test')
26 plt.ylabel('y_test')
27 plt.title('Actual vs. Predicted (Polynomial Regression)')
28 plt.legend()
29
30 # Actual vs. Predicted for ANN
31 plt.subplot(2, 2, 3)
32 plt.scatter(X_test, y_test, color='blue', label='Actual')
33 plt.scatter(X_test, ann_preds, color='green', label='ANN')
34 plt.xlabel('X_test')
35 plt.ylabel('y_test')
36 plt.title('Actual vs. Predicted (ANN)')
37 plt.legend()
38
39 plt.tight_layout()
40 plt.show()
41

```

11/11 [=====] - 0s 2ms/step



▼ TASK6. Compare RMSE, MAE, MSE for the created models.

```

1  from sklearn.metrics import mean_squared_error, mean_absolute_error
2
3  # Function to calculate RMSE, MAE, and MSE
4  def calculate_metrics(y_true, y_pred):
5      rmse = np.sqrt(mean_squared_error(y_true, y_pred))
6      mae = mean_absolute_error(y_true, y_pred)
7      mse = mean_squared_error(y_true, y_pred)
8      return rmse, mae, mse
9
10 # Calculate metrics for Linear Regression
11 linear_reg_rmse, linear_reg_mae, linear_reg_mse = calculate_metrics(y_test, linear_reg_preds)
12
13 # Calculate metrics for Polynomial Regression models
14 poly_reg_metrics = []
15 for degree, poly_reg_model in poly_reg_models:
16     poly_reg_preds = poly_reg_model.predict(poly_features.transform(X_test))
17     rmse, mae, mse = calculate_metrics(y_test, poly_reg_preds)
18     poly_reg_metrics.append((degree, rmse, mae, mse))
19
20 # Calculate metrics for ANN
21 ann_rmse, ann_mae, ann_mse = calculate_metrics(y_test, ann_preds)
22
23 # Print and compare metrics
24 print(f"Linear Regression - RMSE: {linear_reg_rmse}, MAE: {linear_reg_mae}, MSE: {linear_reg_mse}")
25
26 for degree, rmse, mae, mse in poly_reg_metrics:
27     print(f"Polynomial Regression (Degree {degree}) - RMSE: {rmse}, MAE: {mae}, MSE: {mse}")
28
29 print(f"ANN - RMSE: {ann_rmse}, MAE: {ann_mae}, MSE: {ann_mse}")

```

```

Linear Regression - RMSE: 0.12634578473993902, MAE: 0.09770004663970339, MSE: 0.015963257321551006
Polynomial Regression (Degree 5) - RMSE: 0.12095827506545165, MAE: 0.09411205597383832, MSE: 0.014630904306809463
ANN - RMSE: 0.13814090837163306, MAE: 0.10843921525228352, MSE: 0.01908291056573992

```

▼ TASK 7 RESULT

The Polynomial Regression (Degree 5) model has the lowest RMSE and MSE which means it has the smallest average prediction error and the smallest squared prediction errors, respectively.

The Linear Regression model also performs well but has slightly higher RMSE and MSE compared to the Polynomial Regression (Degree 5).

The ANN model has the highest RMSE and MSE, indicating that its predictions have a larger error compared to the other two models.

Therefore, based on the provided metrics, the Polynomial Regression (Degree 5) model appears to be the best-performing model for this dataset. It has the lowest RMSE and MSE, which suggests that it provides the most accurate predictions among the three models. However, it's essential to consider other factors such as model complexity, training time, and interpretability when choosing the best model for a specific application.

1

▼ DATASET-02 - Electricity load forecasting

Time Series Forecasting Loading of Panama with reference to Weather Parameters & Special Days. There are a total 15 features in this data set.

- 12 features are Numerical continuous values : Weather Parameters
- 03 nos. of parameters contains the details of the Special days (Holidays, Holidays_ID, School days)
- No null values in data sets

Output- To be predicted

Loading of the Panama in 'nat_demand' column

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from warnings import simplefilter
6 from sklearn.linear_model import LinearRegression
7 from tensorflow.keras.models import Sequential, Model
8 from tensorflow.keras.layers import Dense, Input, Dropout
9 from tensorflow.keras.utils import to_categorical, plot_model
10 from sklearn.model_selection import train_test_split
11
12 from sklearn.preprocessing import PolynomialFeatures, StandardScaler
13 from sklearn.metrics import mean_squared_error, mean_absolute_error
14 from sklearn.neural_network import MLPRegressor
15 from math import sqrt
16 import tensorflow as tf
17 from tensorflow import keras
18 from tensorflow.keras import layers
19
```

```
1 df = pd.read_csv("https://raw.githubusercontent.com/Jatansahu/DEEP_LEARNING_ASSIGNMENTS/main/LAB_05/E1
```

▼ Data Understanding

```
1 df
```


9/24/23, 11:38 PM202218061_Jatan_Sahu_DL05.ipynb - Colaboratory

	nat_demand	T2M_toc	QV2M_toc	TQL_toc	W2M_toc	T2M_san	QV2M_san	TQL_san	W2M_san	T2M_dav	QV2M_dav	TQL_dav
datetime												
2015-03-01 01:00:00	970.3450	25.865259	0.018576	0.016174	21.850546	23.482446	0.017272	0.001855	10.328949	22.662134	0.016562	0.096100
2015-03-01 02:00:00	912.1755	25.899255	0.018653	0.016418	22.166944	23.399255	0.017265	0.001327	10.681517	22.578943	0.016509	0.087646
2015-03-01 03:00:00	900.2688	25.937280	0.018768	0.015480	22.454911	23.343530	0.017211	0.001428	10.874924	22.531030	0.016479	0.078735
2015-03-01 04:00:00	889.9538	25.957544	0.018890	0.016273	22.110481	23.238794	0.017128	0.002599	10.518620	22.512231	0.016487	0.068390
2015-03-01 05:00:00	893.6865	25.973840	0.018981	0.017281	21.186089	23.075403	0.017059	0.001729	9.733589	22.481653	0.016456	0.064362
...
2019-12-31 19:00:00	1301.6065	26.635645	0.018421	0.013165	13.184052	25.135645	0.018048	0.064240	3.086798	23.620020	0.016697	0.073425
2019-12-31 20:00:00

This is the hourly data of electricity load from 01 March 2015 to 31 December 2019

Dataset Description (With some google search),

nat_demand: National electricity load

T2M: Temperature at 2 meters

QV2M: Relative humidity at 2 meters

TQL: Liquid precipitation

W2M: Wind speed at 2 meters

And after the underscore is the city

toc: Tocumen, Panama city

san: Santiago city

dav: David city

The rest of variables:

Holiday_ID: Unique identification number integer

holiday: Holiday binary indicator (1=holiday, 0=regular day)

school: School period binary indicator (1=school, 0=vacations)

Data sources provide hourly records. The data composition is the following:

Historical electricity load, available on daily post-dispatch reports, from the grid operator (CND).

Historical weekly forecasts available on weekly pre-dispatch reports, both from CND.

Calendar information related to school periods, from Panama's Ministry of Education.

Calendar information related to holidays, from "When on Earth?" website. Weather variables, such as temperature, relative humidity, precipitation, and wind speed, for three main cities in Panama, from Earthdata.

Information Source --> <https://www.kaggle.com/datasets/saurabhshahane/electricity-load-forecasting>

```
1 df.info()


<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 43775 entries, 2015-03-01 01:00:00 to 2019-12-31 23:00:00
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   nat_demand  43775 non-null  float64
1   T2M_toc     43775 non-null  float64
2   QV2M_toc    43775 non-null  float64
3   TQL_toc     43775 non-null  float64
4   W2M_toc     43775 non-null  float64
5   T2M_san     43775 non-null  float64
6   QV2M_san    43775 non-null  float64
7   TQL_san     43775 non-null  float64
```

```

8  W2M_san      43775 non-null float64
9  T2M_dav      43775 non-null float64
10 QV2M_dav     43775 non-null float64
11 TQL_dav      43775 non-null float64
12 W2M_dav      43775 non-null float64
13 Holiday_ID   43775 non-null int64
14 holiday      43775 non-null int64
15 school       43775 non-null int64
dtypes: float64(13), int64(3)
memory usage: 5.7 MB

```

```
1 df['nat_demand'].plot()
```



```

1 def make_lags(df, lags):
2     ret_df = pd.concat([f'lag_{i}':df.shift(i) for i in range(1,lags+1)},axis=1)
3     columns = [''.join(col).strip() for col in ret_df.columns.values]
4     ret_df.columns = columns
5     return ret_df

```

```
1 df.columns
```

```

Index(['nat_demand', 'T2M_toc', 'QV2M_toc', 'TQL_toc', 'W2M_toc', 'T2M_san',
      'QV2M_san', 'TQL_san', 'W2M_san', 'T2M_dav', 'QV2M_dav', 'TQL_dav',
      'W2M_dav', 'Holiday_ID', 'holiday', 'school'],
      dtype='object')

```

▼ Feature Selection

```

1 import pandas as pd
2
3 # Assuming you have a DataFrame 'df' with your data
4 threshold = 0.5 # Set your desired correlation threshold
5
6 # Calculate the correlation of features with 'nat_demand'
7 correlation_with_target = df.corr()['nat_demand'].abs()
8
9 # Create a mask for features to keep (True) or remove (False)
10 mask = (correlation_with_target >= threshold) | (correlation_with_target.isna())
11
12 # Apply the mask to your DataFrame
13 df_filtered = df.loc[:, mask]

```

In this, we keep only the features with a correlation coefficient(Absolute so it counts negative correlatiob also) greater than or equal to the specified threshold.

```
1 df
```

	nat_demand	T2M_toc	QV2M_toc	TQL_toc	W2M_toc	T2M_san	QV2M_san	TQL_san	W2M_san	T2M_dav	QV2M_dav	TQL_dav
datetime												
2015-03-01 01:00:00	970.3450	25.865259	0.018576	0.016174	21.850546	23.482446	0.017272	0.001855	10.328949	22.662134	0.016562	0.096100
2015-03-01 02:00:00	912.1755	25.899255	0.018653	0.016418	22.166944	23.399255	0.017265	0.001327	10.681517	22.578943	0.016509	0.087646
2015-03-01 03:00:00	900.2688	25.937280	0.018768	0.015480	22.454911	23.343530	0.017211	0.001428	10.874924	22.531030	0.016479	0.078735
2015-03-01 04:00:00	889.9538	25.957544	0.018890	0.016273	22.110481	23.238794	0.017128	0.002599	10.518620	22.512231	0.016487	0.068390
2015-03-01 05:00:00	893.6865	25.973840	0.018981	0.017281	21.186089	23.075403	0.017059	0.001729	9.733589	22.481653	0.016456	0.064362
...
2019-12-31 19:00:00	1301.6065	26.635645	0.018421	0.013165	13.184052	25.135645	0.018048	0.064240	3.086798	23.620020	0.016697	0.073425
2019-12-31 20:00:00	1250.9634	26.495935	0.018162	0.014713	13.443892	24.769373	0.017781	0.058838	3.659980	23.284998	0.016606	0.064362
2019-12-31 21:00:00	1193.6802	26.354456	0.017980	0.013836	13.442195	24.479456	0.017606	0.038086	3.769294	23.041956	0.016492	0.054260
2019-12-31 22:00:00	1130.4575	26.166895	0.017965	0.018486	13.420656	24.112207	0.017393	0.020386	3.872397	22.862207	0.016401	0.055557
2019-12-31 23:00:00	1084.4737	25.976373	0.018072	0.023315	13.749788	23.663873	0.017156	0.019531	4.165276	22.726373	0.016302	0.061371

43775 rows × 16 columns

1 df_filtered

	nat_demand	T2M_toc	T2M_san	T2M_dav
datetime				
2015-03-01 01:00:00	970.3450	25.865259	23.482446	22.662134
2015-03-01 02:00:00	912.1755	25.899255	23.399255	22.578943
2015-03-01 03:00:00	900.2688	25.937280	23.343530	22.531030
2015-03-01 04:00:00	889.9538	25.957544	23.238794	22.512231
2015-03-01 05:00:00	893.6865	25.973840	23.075403	22.481653
...
2019-12-31 19:00:00	1301.6065	26.635645	25.135645	23.620020
2019-12-31 20:00:00	1250.9634	26.495935	24.769373	23.284998
2019-12-31 21:00:00	1193.6802	26.354456	24.479456	23.041956
2019-12-31 22:00:00	1130.4575	26.166895	24.112207	22.862207
2019-12-31 23:00:00	1084.4737	25.976373	23.663873	22.726373

43775 rows × 4 columns

```
1 X = df_filtered.loc[:, df_filtered.columns != 'nat_demand']
2 y = df_filtered['nat_demand']
```

```
1 # X_train = np.array(X_train).reshape(-1, 1)
2 # y_train = np.array(y_train) # Convert y_train to a NumPy array
3 # X_test = np.array(X_test).reshape(-1, 1)
4 # y_test = np.array(y_test) # Convert y_test to a NumPy array
```

1 X

	T2M_toc	T2M_san	T2M_dav
datetime			
2015-03-01 01:00:00	25.865259	23.482446	22.662134
2015-03-01 02:00:00	25.899255	23.399255	22.578943
2015-03-01 03:00:00	25.937280	23.343530	22.531030
2015-03-01 04:00:00	25.957544	23.238794	22.512231
2015-03-01 05:00:00	25.973840	23.075403	22.481653
...
2019-12-31 19:00:00	26.635645	25.135645	23.620020
2019-12-31 20:00:00	26.495935	24.769373	23.284998
2019-12-31 21:00:00	26.354456	24.479456	23.041956
2019-12-31 22:00:00	26.166895	24.112207	22.862207
2019-12-31 23:00:00	25.976373	23.663873	22.726373

43775 rows × 3 columns

```
1 X0 = make_lags(X.iloc[:,:], lags=2)
```

```
1 X0
```

	lag_1_T2M_toc	lag_1_T2M_san	lag_1_T2M_dav	lag_2_T2M_toc	lag_2_T2M_san	lag_2_T2M_dav
datetime						
2015-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	NaN
2015-03-01 02:00:00	25.865259	23.482446	22.662134	NaN	NaN	NaN
2015-03-01 03:00:00	25.899255	23.399255	22.578943	25.865259	23.482446	22.662134
2015-03-01 04:00:00	25.937280	23.343530	22.531030	25.899255	23.399255	22.578943
2015-03-01 05:00:00	25.957544	23.238794	22.512231	25.937280	23.343530	22.531030
...
2019-12-31 19:00:00	26.999292	25.733667	24.132104	28.112024	26.893274	24.916711
2019-12-31 20:00:00	26.635645	25.135645	23.620020	26.999292	25.733667	24.132104
2019-12-31 21:00:00	26.495935	24.769373	23.284998	26.635645	25.135645	23.620020
2019-12-31 22:00:00	26.354456	24.479456	23.041956	26.495935	24.769373	23.284998
2019-12-31 23:00:00	26.166895	24.112207	22.862207	26.354456	24.479456	23.041956

43775 rows × 6 columns

```
1 X = pd.concat([X.iloc[:,0:], X0], axis=1)
```

```
1 Xx = X.copy()
```

```
1 X
```

```

T2M_toc    T2M_san    T2M_dav    lag_1_T2M_toc    lag_1_T2M_san    lag_1_T2M_dav    lag_2_T2M_toc    lag_2_T2M_san    lag_2_T2M_dav

datetime
1 X.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 43775 entries, 2015-03-01 01:00:00 to 2019-12-31 23:00:00
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   T2M_toc                43775 non-null  float64
 1   T2M_san                43775 non-null  float64
 2   T2M_dav                43775 non-null  float64
 3   lag_1_T2M_toc          43774 non-null  float64
 4   lag_1_T2M_san          43774 non-null  float64
 5   lag_1_T2M_dav          43774 non-null  float64
 6   lag_2_T2M_toc          43773 non-null  float64
 7   lag_2_T2M_san          43773 non-null  float64
 8   lag_2_T2M_dav          43773 non-null  float64
dtypes: float64(9)
memory usage: 4.3 MB

```

```

1 y

datetime
2015-03-01 01:00:00    970.3450
2015-03-01 02:00:00    912.1755
2015-03-01 03:00:00    900.2688
2015-03-01 04:00:00    889.9538
2015-03-01 05:00:00    893.6865
...
2019-12-31 19:00:00    1301.6065
2019-12-31 20:00:00    1250.9634
2019-12-31 21:00:00    1193.6802
2019-12-31 22:00:00    1130.4575
2019-12-31 23:00:00    1084.4737
Name: nat_demand, Length: 43775, dtype: float64

```

```

1 X.dropna(inplace=True)
2 y, X = y.align(X, join='inner')
3 X = np.asarray(X).astype('float32')

```

```

1 X.shape, y.shape

((43773, 9), (43773,))

```

▼ Splitting dataset

```

1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(X, y,
3                                                    test_size=0.2, shuffle=False)

```

▼ Standardization

```

1 from sklearn.preprocessing import StandardScaler
2 scaler = StandardScaler()
3 X_train = scaler.fit_transform(X_train)
4 X_test = scaler.transform(X_test)

```

```

1 X_train.shape, y_train.shape

((35018, 9), (35018,))

```

```

1 # Reshape data to match model input requirements
2 X_train = np.array(X_train).reshape(-1, 1)
3 y_train = np.array(y_train) # Convert y_train to a NumPy array
4 X_test = np.array(X_test).reshape(-1, 1)
5 y_test = np.array(y_test) # Convert y_test to a NumPy array

```

▼ Model Selection

```

1 #Checking which degree is working best
2 degree = [2,3,4,5]
3 for i in degree:
4     poly = PolynomialFeatures(i)
5     new_x = poly.fit_transform(X_train)
6
7     model = LinearRegression()
8     model.fit(new_x, y_train)
9
10    new_x = poly.fit_transform(X_test)
11
12    y_pred = pd.Series(model.predict(new_x))
13    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
14    print(f"RMSE for degree {i} is", rmse)

```

```

RMSE for degree 2 is 141.24272565501363
RMSE for degree 3 is 148.32982052076815
RMSE for degree 4 is 138.30065491038113
RMSE for degree 5 is 148.74826423043328

```

```

1 # Create and train Linear Regression model
2 linear_reg_model = LinearRegression()
3 linear_reg_model.fit(X_train, y_train)
4
5 # Create and train Polynomial Regression models with different degrees
6 poly_degrees = [4] # Change degree which works best
7 poly_reg_models = []
8
9
10 for degree in poly_degrees:
11     poly_features = PolynomialFeatures(degree=degree)
12     X_train_poly = poly_features.fit_transform(X_train)
13     X_test_poly = poly_features.transform(X_test)
14     poly_reg_model = LinearRegression()
15     poly_reg_model.fit(X_train_poly, y_train)
16
17     poly_reg_models.append((degree, poly_reg_model))
18
19 # Create and train an Artificial Neural Network (ANN) model
20 ann_model = keras.Sequential([
21     layers.Dense(16, activation='relu', input_shape=(9,)),
22     layers.Dense(32, activation='relu'),
23     layers.Dense(64, activation='relu'),
24     layers.Dense(32, activation='relu'),
25     layers.Dense(16, activation='relu'),
26     layers.Dense(1) # Output layer
27 ])
28
29 # Compile the ANN model
30 ann_model.compile(optimizer='adam', loss='mean_squared_error')
31
32 # Train the ANN model
33 ann_model.fit(X_train, y_train, epochs=40, batch_size=32, validation_split=0.2)
34

```

```

Epoch 1/40
876/876 [=====] - 3s 3ms/step - loss: 209721.8594 - val_loss: 21218.4922
Epoch 2/40
876/876 [=====] - 2s 3ms/step - loss: 16871.1797 - val_loss: 18001.6797
Epoch 3/40
876/876 [=====] - 3s 3ms/step - loss: 16251.5732 - val_loss: 18242.0645
Epoch 4/40
876/876 [=====] - 2s 2ms/step - loss: 15881.6230 - val_loss: 15874.8711
Epoch 5/40
876/876 [=====] - 2s 2ms/step - loss: 15592.6826 - val_loss: 14744.7490
Epoch 6/40
876/876 [=====] - 2s 2ms/step - loss: 15588.8125 - val_loss: 19204.2480
Epoch 7/40
876/876 [=====] - 2s 2ms/step - loss: 15451.8330 - val_loss: 14890.7461
Epoch 8/40
876/876 [=====] - 2s 2ms/step - loss: 15377.7949 - val_loss: 14241.8467
Epoch 9/40

```

```
876/876 [=====] - 3s 3ms/step - loss: 15273.3652 - val_loss: 19738.2891
Epoch 10/40
876/876 [=====] - 2s 2ms/step - loss: 15272.5342 - val_loss: 17489.2148
Epoch 11/40
876/876 [=====] - 2s 2ms/step - loss: 15247.3887 - val_loss: 19448.5039
Epoch 12/40
876/876 [=====] - 2s 2ms/step - loss: 15049.1895 - val_loss: 18737.1543
Epoch 13/40
876/876 [=====] - 2s 2ms/step - loss: 15047.5605 - val_loss: 13587.7236
Epoch 14/40
876/876 [=====] - 2s 2ms/step - loss: 14793.4521 - val_loss: 16121.6875
Epoch 15/40
876/876 [=====] - 3s 4ms/step - loss: 14532.1934 - val_loss: 14170.2285
Epoch 16/40
876/876 [=====] - 2s 2ms/step - loss: 14417.3594 - val_loss: 19749.9629
Epoch 17/40
876/876 [=====] - 2s 2ms/step - loss: 14256.9385 - val_loss: 18700.4160
Epoch 18/40
876/876 [=====] - 2s 2ms/step - loss: 14211.5947 - val_loss: 11815.4961
Epoch 19/40
876/876 [=====] - 2s 2ms/step - loss: 13969.0381 - val_loss: 17206.2520
Epoch 20/40
876/876 [=====] - 2s 2ms/step - loss: 13885.5430 - val_loss: 14041.0488
Epoch 21/40
876/876 [=====] - 3s 3ms/step - loss: 13817.2852 - val_loss: 12591.3008
Epoch 22/40
876/876 [=====] - 2s 2ms/step - loss: 13798.7529 - val_loss: 19565.0254
Epoch 23/40
876/876 [=====] - 2s 2ms/step - loss: 13761.7139 - val_loss: 18036.4316
Epoch 24/40
876/876 [=====] - 2s 2ms/step - loss: 13700.1582 - val_loss: 18365.7129
Epoch 25/40
876/876 [=====] - 2s 2ms/step - loss: 13659.7314 - val_loss: 17233.2051
Epoch 26/40
876/876 [=====] - 2s 2ms/step - loss: 13646.6357 - val_loss: 16227.2656
Epoch 27/40
876/876 [=====] - 3s 3ms/step - loss: 13547.4580 - val_loss: 15636.2568
Epoch 28/40
876/876 [=====] - 2s 2ms/step - loss: 13700.1777 - val_loss: 14385.8896
Epoch 29/40
876/876 [=====] - 2s 2ms/step - loss: 13543.9697 - val_loss: 13434.6475
```

```
1 # Predictions for all models
```

```
2 linear_reg_preds = linear_reg_model.predict(X_test)
```

```
3 poly_reg_preds = [poly_reg_model.predict(poly_features.transform(X_test)) for _, poly_reg_model in pol
```

```
4 ann_preds = ann_model.predict(X_test)
```

```
274/274 [=====] - 0s 1ms/step
```

```

1 from sklearn.metrics import mean_squared_error, mean_absolute_error
2
3 # Function to calculate RMSE, MAE, and MSE
4 def calculate_metrics(v_true, v_pred):

```

After adding two more layer and increasing nodes in ANN I found that ANN is performing better than Polynomial regression. Initially I found that Polynomial with degree 4 works best.

Challenge - I am facing challenge while plotting time series graphs. I tried by different methods but not getting better visualization.

```

10 # Calculate metrics for Linear Regression
11 # fig, ax = plt.subplots(figsize=(15, 5))
12 # fig.autofmt_xdate()
13 # plt.plot(X_test, y_test,
14 #          color='darkgray', alpha=0.8, label='Original')
15 # plt.plot(linear_reg_preds, label='Training prediction')
16 # plt.plot(X_test, linear_reg_preds, label='Testing prediction')
17 # plt.legend()
18 # ax.set_xticks([i for i in list(elec_df['datetime'])[::1000]])
19 # plt.xticks(rotation=60)
20 # plt.show()
21 ann_rmse, ann_mae, ann_mse = calculate_metrics(y_test, ann_preds)
22
23
24 print(f"Linear Regression - RMSE: {linear_reg_rmse}, MAE: {linear_reg_mae}, MSE: {linear_reg_mse}")
25
26 for degree, rmse, mae, mse in poly_reg_metrics:
27     print(f"Polynomial Regression (Degree {degree}) - RMSE: {rmse}, MAE: {mae}, MSE: {mse}")
28
29 print(f"ANN - RMSE: {ann_rmse}, MAE: {ann_mae}, MSE: {ann_mse}")

```

Linear Regression - RMSE: 146.6986571506399, MAE: 118.55872228666209, MSE: 21520.496009800987
 Polynomial Regression (Degree 4) - RMSE: 138.30065491038113, MAE: 110.17540567291904, MSE: 19127.071148640323
 ANN - RMSE: 132.0439735506955, MAE: 107.29326549309992, MSE: 17435.610951056773